



PARAM Brahma

User's Manual

Ver. 1.0

Last updated: April 30, 2020

www.cdac.in

Copyright Notice

Copyright © 2020 Centre for Development of Advanced Computing
All Rights Reserved.

Any technical documentation that is made available by C-DAC (Centre for Development of Advanced Computing) is the copyrighted work of C-DAC and is owned by C-DAC. This technical documentation is being delivered to you as is, and C-DAC makes no warranty as to its accuracy or use. Any use of the technical documentation or the information contained therein is at the risk of the user. C-DAC reserves the right to make changes without prior notice.

No part of this publication may be copied without the express written permission of C-DAC.

Trademarks

CDAC, CDAC logo, NSM logo are trademarks or registered trademarks.

Other brands and product names mentioned in this manual may be trademarks or registered trademarks of their respective companies and are hereby acknowledged.



Intended Audience

This document is meant for PARAM Brahma users.

Typographic Conventions

Symbol	Meaning
<u>Blue underlined text</u>	A hyperlink or link you can click to go to a related section in this document or to a URL in your web browser.
Bold	The names of menus, menu items, headings, and buttons.
<i>Italics</i>	Variables or placeholders or special terms in the document.
<i>Console text</i>	Console commands



Getting help

For technical assistance or license renewal, please send an email to brahmasupport@iiserpune.ac.in.

Give us your feedback

We value your feedback. Kindly send your comments on content of this document to brahmasupport@iiserpune.ac.in . Please include the page number of the document along with your feedback.

**DISCLAIMER**

The information contained in this document is subject to change without notice. C-DAC shall not be liable for errors contained herein or for incidental or consequential damages in connection with the performance or use of this manual.

Table of Contents

Introduction.....	7
System Architecture and Configuration.....	8
System Hardware Specifications	8
Master Nodes	8
Login Nodes	8
Service Nodes	9
CPU Compute Nodes	9
Operating System	10
Primary Interconnection Network	11
Secondary Interconnection Network	11
Software Stack.....	11
First Things First.....	14
Getting an Account on PARAM Brahma.....	14
First login	15
Forgot Password?	15
System Access	16
Remote Access	16
Transferring files between local machine and HPC cluster.....	17
Tools	19
Running Interactive Jobs.....	22
Managing Jobs through its Lifecycle.....	23
List Partition	23
Submit the job	24
List jobs.....	26
Get job details	26
Suspend a job	28
Resume a job	28
Kill a job	28
Hold a job	29
Release a job.....	29
Addressing Basic Security Concerns	30

More about Batch Jobs (SLURM)	31
Parameters used in SLURM job script	31
I am familiar with PBS/ TORQUE. How do I migrate to SLURM?	34
Preparing your own Executable	35
Job Scheduling on PARAM Brahma	39
Scheduler	39
sinfo	39
Walltime	40
Scheduling Type.....	41
Debugging Your Codes	44
Introduction.....	44
Basics How Tos	44
Conclusions.....	63
Points to Note.....	63
Overall Coding Modifications Done	64
Machine Learning / Deep Learning Application Development.....	65
How to Install your own Software?	66
Some Important Facts	68
About File Size	68
Little-Endian and Big-Endian issues?	69
Best Practices for HPC	70
Installed Applications/Libraries	71
Standard Application Programs on PARAM Brahma.....	71
LAMMPS Applications	72
OpenFOAM.....	75
WRF Application	77
NAMD Application.....	78
GROMACS Application.....	82
Acknowledging the National Supercomputing Mission in Publications.....	85
Getting Help – PARAM Brahma Support	86

Steps to Create a New Ticket86

Closing Your Account on PARAM Brahma.....88

PARAM Brahma Account Request Form..... 90

References..... 94

List of Figures

Figure 1 - PARAM Brahma Architecture Diagram.....	10
Figure 2 – Software Stack.....	12
Figure 3 - A snapshot of command using MobaXterm	19
Figure 4 - A snapshot of "scp" command using Windows command prompt.....	20
Figure 5 - A snapshot of "scp" command using Windows PowerShell.....	20
Figure 6 - A snapshot of "scp" tool to transfer file to and from remote computer.	21
Figure 7 – Output of sinfo command.....	24
Figure 8 – snapshot depicting the usage of “Job Array”	25
Figure 9 – View of available partition and node information on the cluster	39
Figure 10 - Listing the shares of association to a cluster	42
Figure 11 – snapshot of debugging process	47
Figure 12 – Snapshot of debugging process	48
Figure 13- output at a debugging stage.....	49
Figure 14 – Snapshot of debugging process	49
Figure 15 – Output depicting “Arithmetic Exception”	50
Figure 16 – Snapshot of debugging process	50
Figure 17 – Well, we dumped core !!	51
Figure 18 - Snapshot of debugging process.....	51
Figure 19 – setting breakpoint.....	52
Figure 20 – single stepping through to catch error !!.....	53
Figure 21 – Debugging continued	54
Figure 22 – Debugging continued	54
Figure 23 – Setting a watch point	55
Figure 24 – Debugging continued	56
Figure 25 – Well, Back to square one !!.....	57
Figure 26 – Again Dumping Core!! Things are getting interesting or frustrating or both!!	57
Figure 27 – Debugging continued	58
Figure 28 – Debugging continued	58
Figure 29 – Debugging continued (Will it ever end?)	59
Figure 30 – We are almost there !!.....	59
Figure 31 – Debugging continued	60
Figure 32 – At last a clue!!!	61
Figure 33 - Correction applied !!.....	62
Figure 34 – Resolved !!!	63
Figure 35 – What all we did to get things right !	64
Figure 36 – Snapshot of Ticketing System	86
Figure 37 - Snapshot of Ticketing System.....	87
Figure 38 - Snapshot of Ticketing System.....	87
Figure 39 - Snapshot of Ticketing System.....	88

Introduction

This document is the user manual for the PARAM Brahma Supercomputing facility at IISER-Pune. It covers a wide range of topics ranging from a detailed description of the hardware infrastructure to the information required to utilize the supercomputer, such as information about logging on to the supercomputer, submitting jobs, retrieving the results on to user's Laptop/ Desktop etc. In short, the manual describes all that one needs know to effectively utilize PARAM Brahma.

The supercomputer PARAM Brahma is based on Intel Xeon Cascade lake processors. The system was designed and implemented by HPC Technologies team, Centre for Development of Advanced Computing (C-DAC).

It consists of 2 Master nodes, 4 Login nodes, 4 Service nodes and 179 compute nodes with total peak computing capacity of **797 TFLOPS** performance.

System Architecture and Configuration

System Hardware Specifications

PARAM Brahma system is based on processor Intel Xeon Platinum 8268 with total peak performance of 797 TFLOPS. The cluster consists of compute nodes connected with BullSequana XH2000 HDR 100 infiniband interconnect network. The system uses the Lustre parallel file system.

- Total number of nodes: 189 (10 + 179)
 - Master nodes: 2
 - Login nodes: 4
 - Service nodes: 4
 - CPU nodes: 179

Master Nodes

PARAM Brahma is an aggregation of a large number of computers connected through networks. The basic purpose of the master node is to manage and monitor each of the constituent component of PARAM Brahma from a system’s perspective. This involves operations like monitoring the health of the components, the load on the components, the utilization of various sub-components of the computers in PARAM Brahma.

Master Nodes : 2	
2* Intel Xeon G-6248	Total Cores = 80 cores
Cores =40, 2.5 GHz	
Memory= 384 GB	Total Memory = 768 GB
HDD = 1.2 TB x 8	

Login Nodes

Login nodes are typically used for administrative tasks such as editing, writing scripts, transferring files, managing your jobs and the like. You will always get connected to one of the login nodes. From the login nodes you can get connected to a compute node and execute and interactive job or submit batch jobs through the batch system (SLURM) to run your jobs on compute nodes. For ALL users PARAM Brahma login nodes are the entry points

and hence are shared. By default, there will be a limit on the CPU time that can be used on a login node by a user and there is a limit/user on the memory as well. If any of these are exceeded, the job will get terminated.

Login Nodes : 4

2* Intel Xeon G-6248

Cores = 40, 2.5 GHz

Memory= 384 GB

HDD = 1.2 TB x 8

Total Cores = 160 cores

Total Memory = 1536 GB

Service Nodes

Typically, the purpose of the service node is to provide Security, Management, monitoring and other services to the cluster.

Service Nodes : 4

2* Intel Xeon G-6248

Cores = 40, 2.5 GHz

Memory= 384 GB

HDD = 1.2 TB x 8

Total Cores = 160 cores

Total Memory= 1536 GB

CPU Compute Nodes

CPU nodes are indeed the work horses of PARAM Brahma. All the CPU intensive activities are carried on these nodes. Users can access these nodes from the login node to run interactive or batch jobs. Some of the nodes have higher memory, which can be exploited by users in the aforementioned way.

CPU only Compute Nodes : 140

2* Intel Xeon Platinum 8268

Cores = 48, 2.9 GHz

Memory= 192 GB, DDR4 2933 MHz

SSD = 480 GB (local scratch) per node

Total Cores = 6720 cores

Total Memory=26880 GB

CPU only Compute Nodes with High memory : 39

2* Intel Xeon Platinum 8268

Cores = 48, 2.9 GHz

Memory= 768 GB, DDR4 2933 MHz

SSD = 480 GB (local scratch) per node

Total Cores = 1872 cores

Total Memory=29952 GB

Storage

- Based on Lustre parallel file system
- Total useable capacity 800 TiB Primary storage and Archival storage 250 TiB
- Throughput 25GB/s

Operating System

- Operating system on PARAM Brahma is Linux – CentOS 7.6

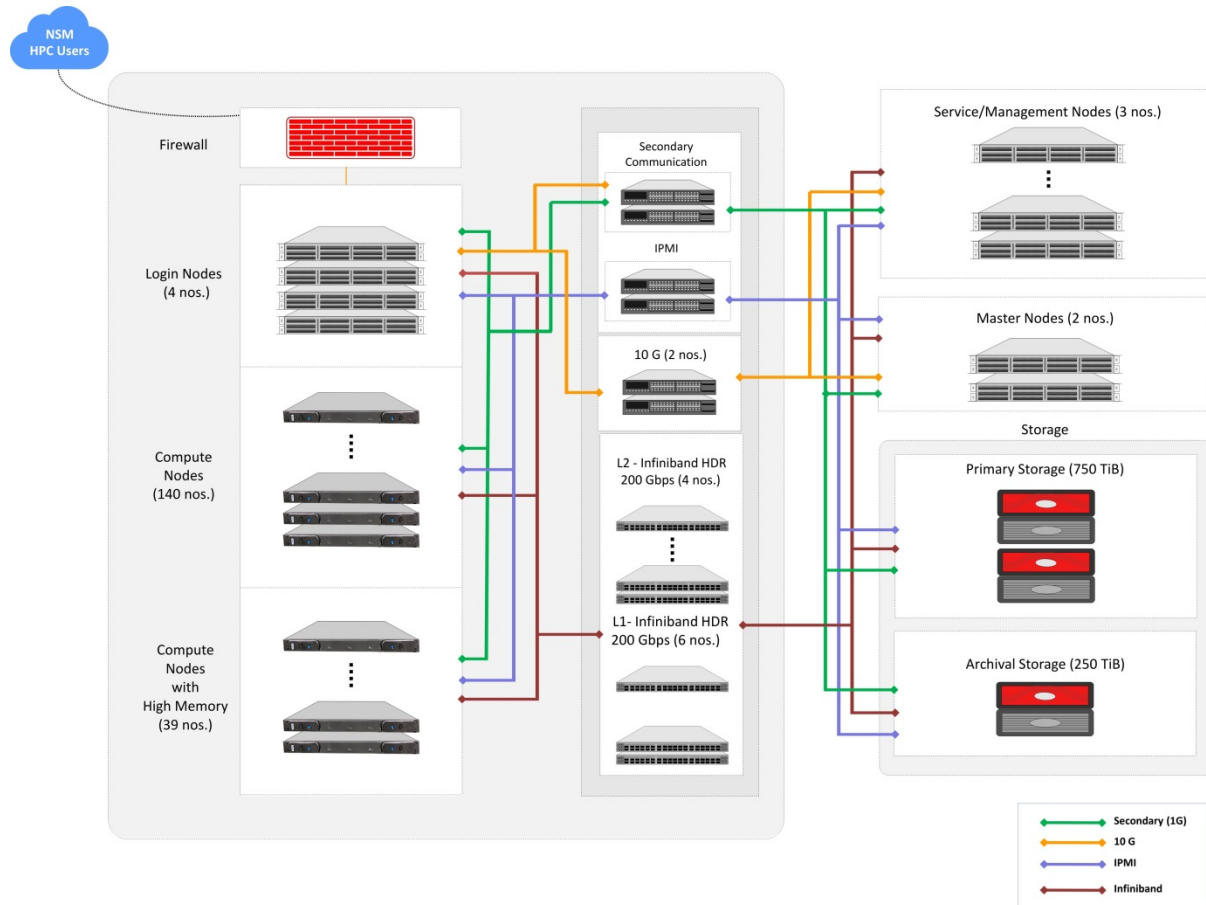


Figure 1 - PARAM Brahma Architecture Diagram

Network infrastructure

A robust network infrastructure is essential to implement the basic functionalities of a cluster. These functionalities are:

- Management functionalities i.e. to monitor, trouble shoot, start, stop various components of the cluster, etc. (Network/ portion of Network which implements this functionality is referred to as Management fabric).
- Ensuring fast read/ write access to the storage (Network/ portion of Network which implements this functionality is referred to as storage fabric).

- c) Ensuring fast I/O operations like connecting to other clusters, connecting the cluster to various users on the campus LAN, etc. (Network/ portion of Network which implements this functionality is referred to as I/O Fabric).
- d) Ensuring High-Bandwidth, Low-latency communication amongst processors to for achieving high-scalability (Network/ portion of Network which implements this functionality is referred to as Message Passing Fabric)

Technically, ALL the aforementioned functionalities can be implemented in a single network. From the perspectives of requirements, optimal performance and economic suitability, the aforementioned functionalities are implemented using two different networks based on different technologies, as mentioned next:

Primary Interconnection Network

Computing nodes of PARAM Brahma are interconnected by a high-bandwidth, low-latency interconnect network.

InfiniBand: HDR 100 Gbps

InfiniBand is a high-performance communication architecture owned by Mellanox. This communication architecture offers low communication latency, low power consumption and a high throughput. All CPU nodes are connected via InfiniBand interconnect network.

Secondary Interconnection Network

Gigabit Ethernet: 1 Gbps

Gigabit Ethernet is the interconnection network that is most commonly available. For Gigabit Ethernet, no additional modules or libraries are required. The Open MPI, MPICH implementations will work over Gigabit Ethernet.

Software Stack

Software stack is an aggregation of software components that work in tandem to accomplish a given task. The task can be, to facilitate a user to execute his job/s or to facilitate a system administrator to manage a system efficiently. In effect, the software will have all the necessary components to accomplish a given task. There may be multiple components of different flavors to accomplish a given sub-task. The user/ administrator may mix and match these components depending on his choice. Typically, a user would be interested in preparing his executable, executing the same with his data sets and visualize the output generated by him. For accomplishing the same, the user would need to compile his codes, link the codes with communication libraries, Math Libraries, Numerical algorithm

libraries, prepare the executables, run the same with desired data sets, monitor the progress of his jobs, gathering the results and visualizing the output.

Typically, a system administrator would be interested in ensuring that all the resources are optimally utilized. For accomplishing this, he may need some installation tools, tools for checking the health of all the components, good schedulers, tools to facilitate allocation of resources to users and monitor the usage of the resources.

The software stack provided with this system has a gamut of software components which meets all the requirements of a user and that of a system administrator. The components of the software stack are depicted in figure 2.

Amongst these, C-CHAKSHU and CHReME have been recently developed and deployed by CDAC. We solicit your feedback on these tools at brahmasupport@iiserpune.ac.in

C-CHAKSHU: This is a multi-cluster management tool that facilitates the administrator to efficiently operate the HPC facility. It also enables the user to monitor system metrics relating to CPU, Storage, Interconnects, File system and Application specific utilization from a single dashboard. For more information, please follow the link <http://parambrahma.iiserpune.ac.in:4200>.

CHReME: This is a web-based resource management portal with an intuitive GUI that facilitates managing, monitoring and tracking of HPC workloads. It provides a customizable and user-friendly workflow interface to HPC users for state-based pre-processing, execution and post-processing of applications. For more details please follow the link <http://parambrahma.iiserpune.ac.in:8097/CHReME>.

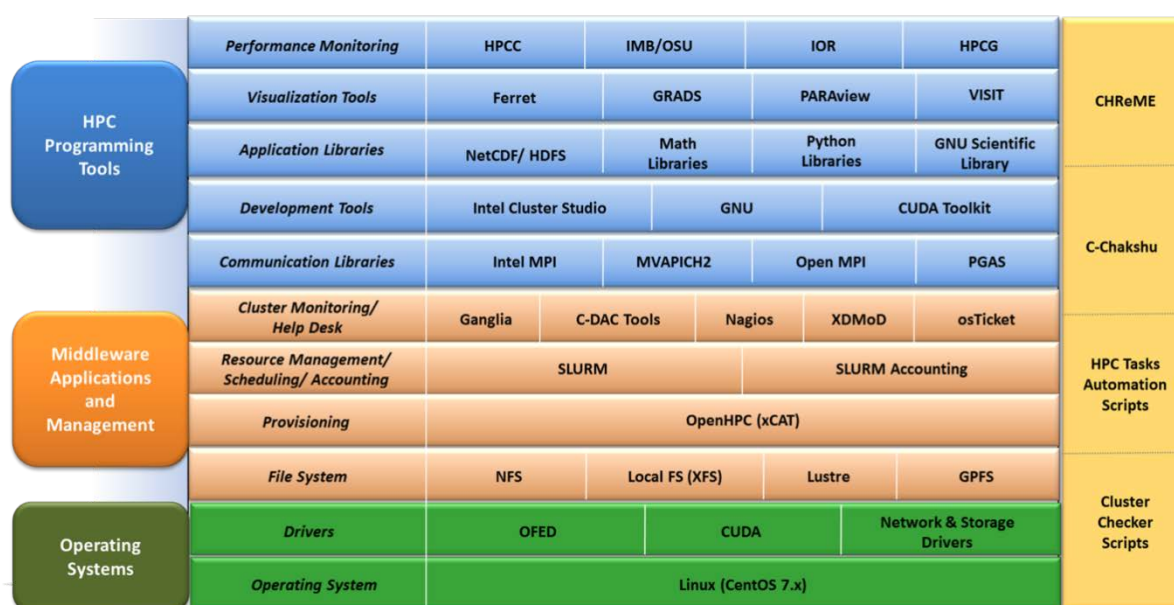


Figure 2 – Software Stack

Functional Areas	Components
Base OS	CentOS 7.6
Architecture	X86_64
Provisioning	xCAT 2.14.6
Cluster Manager	Openhpc (ohpc-xCAT 1.3.8)
Monitoring Tools	C-CHAKSHU, Nagios, Ganglia, XDMoD
Resource Manager	Slurm
I/O Services	Lustre Client
High Speed Interconnects	Mellanox InfiniBand
Compiler Families	GNU (gcc, g++, gfortran) Intel Compiler (icc, ifort, icpc)
MPI Families	MVAPICH, OpenMPI, MPICH

First Things First

Getting an Account on PARAM Brahma

To begin with, you need to get an account on PARAM Brahma. This is a very easy process. Please follow the steps given below:

- a) Download the 'User Account Creation Form' by following the link
- b) <https://parambrahma.iiserpune.ac.in/ucform>
- c) Fill the relevant details.
- d) Get the signatures of your Head of the Department and the 'approving authority'.

Note:

- For IISER users will have the approving authority from IISER Pune. They can submit it to the PARAM Brahma system Administrator* or email a scanned copy to brahmasupport@iiserpune.ac.in.
 - For Users who are not from IISER, Pune will have to email the scanned copy to brahmasupport@iiserpune.ac.in , HoD HPC-Tech CDAC will be the approving authority for them.
- e) You will receive an Email in your official Email ID intimating the creation of your account along with a temporary password set by the system to your account. You will also get a copy of this document by Email.
 - f) Log into PARAM Brahma and you will be prompted to change the password. When once you change the temporary password provide by the system to your own password, you are ready to use PARAM Brahma!!

Info: *

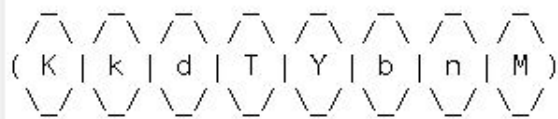
Param Brahma, Indian Institute of Science Education and Research (IISER)
Dr. Homi Bhabha Road,
Pashan, Pune 411 008.
INDIA Email: brahmasupport@iiserpune.ac.in

First login

Whenever the newly created user on PARAM Brahma tries to login with the user Id and password (temporary, system generated) provided over the Email through PARAM Brahma support, he/she will next be prompted to create a “new password” of their choice which will change the temporary, system generated password. This will enable you to keep your account secure. It is recommended that you have a strong password which contains the combination of alphabets (lower case / upper case), numbers, and a few special characters that you can easily remember.

Given next is a screenshot that describes the scenario for “first login”

Observe the picture below and answer the question listed afterwards:



(K | k | d | T | Y | b | n | M)

Type the string above: KkdTYbnM

Password:

You are required to change your password immediately (password aged)
password expired 18078 days ago

New password:

Your password will be valid for 90 days. On expiry of 90 days period, you will be prompted to change your password, on attempting to log in. You are required to provide a new password.

Forgot Password?

There is nothing to panic!! Please raise a ticket regarding this issue and the system administrators will resolve your problem. Please refer to the section “Getting Help – PARAM Brahma Support, described elsewhere in this manual. Follow the GUI based, user-friendly ticketing system. Please follow the steps given below:

1. Open the PARAM Brahma support site i.e the ticketing tool by following the link <https://parambrahma.iiserpune.ac.in/support>
2. Login with your registered email id, Complete name, Contact number.
3. There you can raise a ticket to get the password reset.
4. The system admin person will revert with an email for verification.
5. Once acknowledged, the password is reset for the user and an email is sent back for intimating same.
6. Then the user can login with the temporary password and can set a new password of his/her choice.

System Access

Accessing the cluster

The cluster can be accessed through 4 general login nodes, which allows users to login.

- You may access login node through ssh.
- The login node is primary gateway to the rest of the cluster, which has a job scheduler (called Slurm). You may submit jobs to the queue and they will run when the required resources are available.
- Please do not run programs directly on login node. Login node is use to submit jobs, transfer data and to compile source code. (If your compilation takes more than a few minutes, you should submit the compilation job into the queue to be run on the cluster.)
- By default, two directories are available (i.e. /home and /scratch). These directories are available on login node as well as the other nodes on the cluster. /scratch is for temporary data storage, generally used to store data required for running jobs.

Remote Access

Using SSH in Windows

To access PARAM Brahma you need to “ssh” the login server. PuTTY is the most popular open source “ssh” client application for Windows, you can Download it from (<http://www.putty.org/>). Once installed, find the PuTTY application shortcut in your Start Menu, desktop. On clicking the PuTTY icon The PuTTY Configuration dialog should appear. Locate the “Host Name or IP Address” input Field in the PuTTY Configuration screen. Enter the user name along with IP address or Hostname with which you wish to connect.

(e.g. [username]@parambrahma.iiserpune.ac.in)

Enter your password when prompted, and press Enter.

Using SSH in Mac or Linux

Both Mac and Linux systems provide a built-in SSH client, so there is no need to install any additional package. Open the terminal, connect to an SSH server by typing the following command:

```
ssh [username]@[hostname]
```

For example, to connect to the PARAM Brahma Login Node, with the username

```
user1: ssh user1@parambrahma.iiserpune.ac.in -p 4422(out side access)
```

You will be prompted for a password, and then will be connected to the server.

Password

How to change the user password?

Use the **passwd** command to change the password for the user from login node.

```
[nikhlesh@login1 ~]$ passwd
Changing password for user nikhlesh.
(current) LDAP Password:
New password:
Retype new password:
```

Transferring files between local machine and HPC cluster

Users need to have the data and application related to their project/research work on PARAM Brahma.

To store the data special directories have been made available to the users with name “scratch and home” the path to this directory is “/scratch” and “/home”. Whereas these directories are common to all the users, a user will get his own directory with their username in /scratch/ as well as /home/ directories where they can store their data.

```
/home/<username>/: ! This directory is generally used by the user to
install applications.
```

```
/scratch/<username>/: ! This directory is user to store the user data
related to the project/research.
```

However, there is limit to the storage provided to the users, the limits have been defined according to quota over these directories, all users will be allotted same quota by default. When a user wishes to transfer data from their local system (laptop/desktop) to HPC system, they can use various methods and tools.

A user using ‘Windows’ operating system will get methods and tools that are native to Microsoft windows and tools that could be installed on your Microsoft windows machine. Linux operating system users do not require any tool. They can just use “scp” command on their terminal, as mentioned below.

Users are advised to keep a copy of their data with themselves, once the project/research work is completed by transferring the data in from PARAM Brahma to their local system (laptop/desktop). The command shown below can be used for effecting file transfers (In all the tools):

```
Scp -r <path to the local data directory> <your username>@<IP of parambrahma>:<path to directory on HPC where to save the data>
```

Example:

Same Command could be used to transfer data from HPC system to your local system (laptop/desktop).

```
Scp -r /dir/dir/file saurabh@<cluster IP/Name>:/home/Saurabh
```

Example:

```
Scp -r <path to directory on HPC> <your username>@<IP of local system>:<path to the local data directory>
```

```
Scp -r /home/saurabh saurabh@<local system IP/Name>:/dir/dir/file
```

Note: The Local system (laptop/desktop) should be connected to the network with which it can access the HPC system.

To reiterate,

Copying Directory/File from local machine to PARAM Brahma:

To copy a local directory from your Linux system (say Wrf-2.0) to your home directory in your PARAM Brahma HPC account, the procedure is:

1. From terminal go to the parent directory using cd command.
user1@mylaptop:~\$cd ~/MyData/
2. Under parent directory type ls <& press Enter key>, & notice Wrf-2.0 is there.
user1@mylaptop: ~\$ls Files TempFiles-0.5 Wrf-2.0
3. Begin copy by typing:
user1@mylaptop:~\$ scp -r Wrf-2.0 (username)@parambrahma.iiserpune.ac.in
< you will be prompted for password ; enter your password >

4. Now login to your account as: `user1@mylaptop:~$ ssh (your username)@parambrahma.iiserpune.ac.in` < you will be prompted for password ; enter password
> `[user1@login ~]$`
5. Is command, you should see Wrf-2.0 directory.
6. While copying from PARAM Brahma to your local machine, follow the same steps

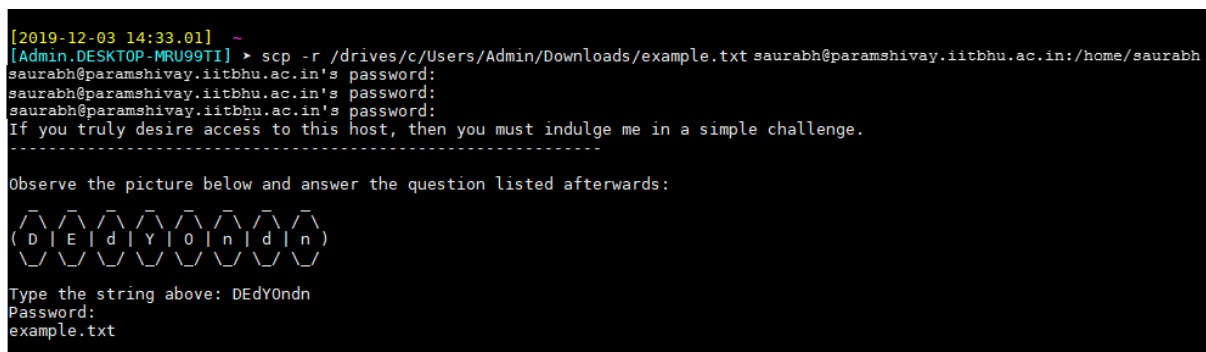
By interchanging source and destination in the scp command. Refer to the generic copying described earlier.

Tools

MobaXterm (Windows installable application):

It is a third party freely available tool which can be used to access the HPC system and transfer file to PARAM Brahma system through your local systems (laptop/desktop).

Link to download this tool : <https://mobaxterm.mobatek.net/download-home-edition.html>



```
[2019-12-03 14:33.01] ~
[Admin.DESKTOP-MRU99TI] > scp -r /drives/c/Users/Admin/Downloads/example.txt saurabh@paramshivay.iitbhu.ac.in:/home/saurabh
saurabh@paramshivay.iitbhu.ac.in's password:
saurabh@paramshivay.iitbhu.ac.in's password:
saurabh@paramshivay.iitbhu.ac.in's password:
If you truly desire access to this host, then you must indulge me in a simple challenge.
-----
Observe the picture below and answer the question listed afterwards:
( D | E | d | Y | O | n | d | n )
Type the string above: DEdYOn dn
Password:
example.txt
```

Figure 3 - A snapshot of command using MobaXterm

Command Prompt (Windows native application):

This is a native tool for Windows machine which can be used to transfer data from PARAM Brahma system through your local systems (laptop/desktop).

```

C:\> Select Command Prompt
Microsoft Windows [Version 10.0.17763.914]
(c) 2018 Microsoft Corporation. All rights reserved.

C:\Users\abhi> scp Desktop/example.txt abhijit@paramshivay.iitbhu.ac.in:/home/abhijit
The authenticity of host 'paramshivay.iitbhu.ac.in (117.220.112.48)' can't be established.
ECDSA key fingerprint is SHA256:GzUVgXJdkmhdqHilVj/GPVC5o1lz+AK3jNTR/h0z4+o.
Are you sure you want to continue connecting (yes/no)?
Warning: Permanently added 'paramshivay.iitbhu.ac.in' (ECDSA) to the list of known hosts.
#####
#                               Welcome to PARAM-Shivay                               #
#                               All connections are monitored and recorded                #
#                               Disconnect IMMEDIATELY if you are not an authorized user!  #
#                               #####                                                    #
If you truly desire access to this host, then you must indulge me in a simple challenge.
-----

Observe the picture below and answer the question listed afterwards:

  / \ / \ / \ / \ / \ / \ / \
 ( C | h | R | h | d | X | d | u )
  \ / \ / \ / \ / \ / \ / \

Type the string above:
Password:
example.txt

```

Figure 4 - A snapshot of "scp" command using Windows command prompt.

PowerShell (Windows native application):

This is a This is a native tool for Windows machine which could be used to transfer data from PARAM Brahma system through your local systems (laptop/desktop).

```

PS C:\Users\abhi> scp Desktop/example.txt abhijit@paramshivay.iitbhu.ac.in:/home/abhijit
The authenticity of host 'paramshivay.iitbhu.ac.in (117.220.112.48)' can't be established.
ECDSA key fingerprint is SHA256:GzUVgXJdkmhdqHilVj/GPVC5o1lz+AK3jNTR/h0z4+o.
Are you sure you want to continue connecting (yes/no)?
Warning: Permanently added 'paramshivay.iitbhu.ac.in,117.220.112.48' (ECDSA) to the list of known hosts.
#####
#                               Welcome to PARAM-Shivay                               #
#                               All connections are monitored and recorded                #
#                               Disconnect IMMEDIATELY if you are not an authorized user!  #
#                               #####                                                    #
If you truly desire access to this host, then you must indulge me in a simple challenge.
-----

Observe the picture below and answer the question listed afterwards:

  / \ / \ / \ / \ / \ / \ / \
 ( G | E | v | M | j | X | m | w )
  \ / \ / \ / \ / \ / \ / \

Type the string above:
Password:
example.txt
PS C:\Users\abhi>

```

Figure 5 - A snapshot of "scp" command using Windows PowerShell.

WinSCP (Windows installable application):

This popular tool is freely available and is used very often to transfer data from Windows machine to Linux machine. This tool is GUI based which makes it very user-friendly.

Link for this tool is : <https://winscp.net/eng/download.php>

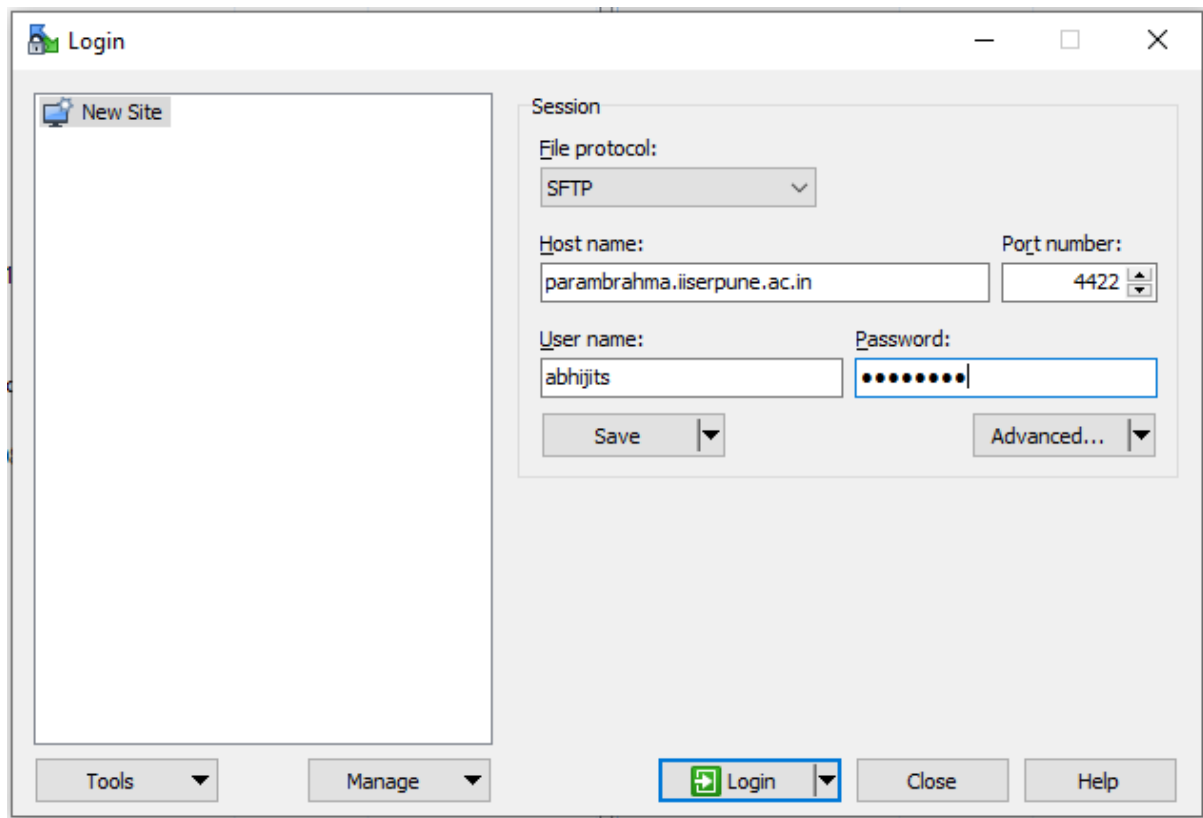


Figure 6 - A snapshot of "scp" tool to transfer file to and from remote computer.

Note:- Port Used for SFTP connection is 4422 and not 22. Please change it to 4422

Running Interactive Jobs

In general, the jobs can be run in an interactive manner or in batch mode. You can run an interactive job as follows:

The following command asks for a single core on one hour with default amount of memory.

```
$ srun --nodes=1 --ntasks-per-node=1 --time=01:00:00 --pty bash -i
```

The command prompt will appear as soon as the job starts. This is how it looks once the interactive job starts:

```
srun: job xxxxx queued and waiting for resources srun: job xxxxx has been
allocated resources
```

Where xxxxx is the job id.

Exit the bash shell to end the job. If you exceed the time or memory limits the job will also abort.

Please note that PARAM Brahma is NOT meant for executing interactive jobs. However, for the purpose of quickly ascertaining successful run of a job before submitting a large job in batch (with large iteration counts), this can be used. This can even be used for running small jobs. The point to be kept in mind is that, since others too would be using this node, it is prudent not to inconvenience them by running large jobs.

It is a good idea to specify the CPU account name as well (if you face any problems)

```
$ srun --account=<NAME_OF_MY_ACCOUNT> --nodes=1 --ntasks-per-node=1 --
time=01:00:00 -- pty bash -i
```

Managing Jobs through its Lifecycle

PARAM Brahma extensively uses modules. The purpose of module is to provide the production environment for a given application, outside of the application itself. This also specifies which version of the application is available for a given session. All applications and libraries are made available through module files. A User has to load the appropriate module from the available modules.

<code>module avail</code>	# This command lists all the available modules
<code>module load intel/2018.5.274</code>	# This will load the intel compilers into your environment
<code>module load intel/2018.5.274</code>	# This will remove all environment setting related to intel-2018 compiler loaded previously

A simple Slurm job script

```
#!/bin/sh
#SBATCH -N 16                                // specifies number of nodes
#SBATCH --ntasks-per-node=40 // specifies core per node
#SBATCH --time=06:50:20 // specifies maximum duration of run
#SBATCH --job-name=lammps                    // specifies job name
#SBATCH --error=job.%J.err_node_40 // specifies error file name
#SBATCH --output=job.%J.out_node_40 //specifies output file name
#SBATCH --partition= standard                // specifies queue name
export I_MPI_FABRICS=shm:dapl
hostname
```

List Partition

`sinfo` displays information about nodes and partitions(queues).

\$ sinfo

```

PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
standard* up   4-00:00:00    3  down* cn[009,102,189]
standard* up   4-00:00:00    3    mix cn[084-086]
standard* up   4-00:00:00   81 alloc cn[001-008,010-020,022,025,027-028,030-035,042-043,046-063,091-092,103-104,107-120,136-145,152-153,186-187]
standard* up   4-00:00:00    4  resv cn[021,024,026,029]
standard* up   4-00:00:00   91  idle cn[023,036-041,044-045,064-083,087-090,093-101,105-106,121-135,146-151,154-172,185,188,190-194]
cpu        up   4-00:00:00    2  down* cn[102,189]
cpu        up   4-00:00:00    3    mix cn[084-086]
cpu        up   4-00:00:00   52 alloc cn[042-043,046-063,091-092,103-104,107-120,136-145,152-153,186-187]
cpu        up   4-00:00:00   86  idle cn[040-041,044-045,064-083,087-090,093-101,105-106,121-135,146-151,154-172,185,188,190-194]
hm         up   4-00:00:00    1  down* cn009
hm         up   4-00:00:00   29 alloc cn[001-008,010-020,022,025,027-028,030-035]
hm         up   4-00:00:00    4  resv cn[021,024,026,029]
hm         up   4-00:00:00    5  idle cn[023,036-039]

```

Figure 7 – Output of sinfo command

Submit the job

We can consider three cases of submitting a job

1. Submitting a simple standalone job

This is a simple submit script which is to be submitted

```

$ sbatch slurm-job.sh
Submitted batch job 106

```

2. Submit a job that's dependent on a prerequisite job being completed

Consider a requirement of pre-processing a job before proceeding to actual processing. Pre-processing is generally done on a single core. In this scenario, the actual processing script is dependent on the outcome of pre-processing script.

here's a simple job script. Note that the Slurm -J option is used to give the job a name.

```

#!/usr/bin/env bash
#SBATCH -p standard
#SBATCH -J simple
sleep 60
Submit the job: $ sbatch simple.sh
Submitted batch job 149

```

Now we'll submit another job that's dependent on the previous job. There are many ways to specify the dependency conditions, but the "singleton" method is the simplest. The Slurm -d singleton argument tells Slurm not to dispatch this job until all previous jobs with the same name have completed.

```

$ sbatch -d singleton simple.sh //may be used for first pre-processing
on a core and then submitting
Submitted batch job 150
$ squeue
  JOBID PARTITION  NAME  USER  ST  TIME  NODES NODELIST(REASON)
    150 standard   simple user1  PD   0:00    1 (Dependency)
    149 standard   simple user1   R   0:17    1 atom01

```

Once the prerequisite job finishes the dependent job is dispatched.

```
$ squeue
JOBID PARTITION NAME USER ST TIME  NODES NODELIST(REASON)
 150 standard  simple user1  R   0:31  1 atom01
```

3. Submit a job with a reservation allocated

Slurm has the ability to reserve resources for jobs being executed. A resource reservation identifies the resources in that reservation and a time period during which the reservation is available. The resources which can be reserved include cores, nodes, licenses and/or burst buffers. When cores or nodes are reserved for a specific user, then user using that reservation can use only those resources and no other user can use those resources.

Reservations can be created, updated, or destroyed only by user root. User can mail the requirement to system admin to create the reservation.

Use the command given below to check the reservation name allocated to your user account

```
$ scontrol show reserv
```

If your ‘user account’ is associated with any reservation the above command will show you the same. For eg. The reservation name given is user_11. Use the command given below to make use of this reservation

```
$ sbatch --reservation=user_11 simple.sh
```

4. Submitting multiple jobs with minor or no changes (array jobs)

A **SLURM job array** is a collection of jobs that differs from each other by only a single index parameter.

```
# Submit a job array with index Values between 0 and 31
$ sbatch -array=0-31 -N1 script.sh
# Submit a job array with index values of 1, 3, 5 and 7
$ sbatch -array=1,3,5,7 -N1 script.sh
# Submit a job array with index values between 1 and 7 with a step size
of 2 (i.e. 1,3,5 and 7)
$ sbatch -array=1-7:2 -N1 script.sh
#Here in above commands -N1 specifies that we are submitting the job on
single node. -N is the parameter to specify number of nodes to be used
to run the job.
```

Figure 8 – snapshot depicting the usage of “Job Array”

Below is the script to submit job array

```
#!/bin/bash
#SBATCH -N 1
#SBATCH --ntasks-per-node=48
#SBATCH --error=job.%A_%a.err
#SBATCH --output=job.%A_%a.out
#SBATCH --time=01:00:00
#SBATCH --partition=standard

module load intel/2018.5.274
cd /home/user1
export OMP_NUM_THREADS=${SLURM_ARRAY_TASK_ID}
/home/user1/md_omp
```

List jobs

Monitoring jobs on SLURM can be done using the command **squeue**. **squeue** is used to view job and job step information for jobs managed by SLURM.

```
$ squeue
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
106 standard slurm-jo user1 R 0:04 1 atom01
```

Get job details

scontrol can be used to report more detailed information about nodes, partitions, jobs, job steps, and configuration.

scontrol show node - shows detailed information about compute nodes.

```
[root@login4 ~]# scontrol show node cn001
NodeName=cn001 Arch=x86_64 CoresPerSocket=24
CPUAlloc=0 CPUTot=48 CPULoad=0.01
AvailableFeatures=(null)
ActiveFeatures=(null)
Gres=(null)
NodeAddr=cn001 NodeHostName=cn001 Version=19.05.2
OS=Linux 3.10.0-957.el7.x86_64 #1 SMP Thu Nov 8 23:39:32 UTC 2018
RealMemory=772828 AllocMem=0 FreeMem=764531 Sockets=2 Boards=1
State=RESERVED ThreadsPerCore=1 TmpDisk=386414 Weight=1 Owner=N/A
MCS_label=N/A
Partitions=standard,hm
BootTime=2020-03-27T01:20:57 SlurmdStartTime=2020-04-13T00:31:05
CfgTRES=cpu=48,mem=772828M,billing=236
AllocTRES=
CapWatts=n/a
CurrentWatts=0 AveWatts=0
ExtSensorsJoules=n/s ExtSensorsWatts=0 ExtSensorsTemp=n/s
```

scontrol show partition - shows detailed information about a specific partition

```
[root@login4 ~]# scontrol show partition standard
PartitionName=standard
  AllowGroups=ALL AllowAccounts=ALL AllowQos=ALL
  AllocNodes=ALL Default=YES QoS=N/A
  DefaultTime=NONE DisableRootJobs=NO ExclusiveUser=NO GraceTime=0
Hidden=NO
  MaxNodes=UNLIMITED MaxTime=4-00:00:00 MinNodes=1 LLN=NO
MaxCPUsPerNode=UNLIMITED
  Nodes=cn[001-136],cn[185-194],cn[137-172]
  PriorityJobFactor=1 PriorityTier=1 RootOnly=NO ReqResv=NO
OverSubscribe=NO
  OverTimeLimit=NONE PreemptMode=OFF
  State=UP TotalCPUs=8736 TotalNodes=182 SelectTypeParameters=NONE
  JobDefaults=(null)
  DefMemPerNode=UNLIMITED MaxMemPerNode=UNLIMITED
  TRESBillingWeights=CPU=1.0,Mem=0.25G,GRES/gpu=10.0
```

scontrol show job - shows detailed information about a specific job or all jobs if no job id is given.

```
$ scontrol show job 106
JobId=106 JobName=test_run
  UserId=user1(7006) GroupId=user1(7006) MCS_label=N/A
  Priority=27111 Nice=0 Account=user1 QOS=normal
  JobState=RUNNING Reason=None Dependency=(null)
  Requeue=1 Restarts=0 BatchFlag=1 Reboot=0 ExitCode=0:0
  RunTime=20:39:18 TimeLimit=4-00:00:00 TimeMin=N/A
  SubmitTime=2020-04-19T14:10:39 EligibleTime=2020-04-19T14:10:39
  AccrueTime=2020-04-19T14:10:39
  StartTime=2020-04-19T14:10:59 EndTime=2020-04-23T14:10:59 Deadline=N/A
  SuspendTime=None SecsPreSuspend=0 LastSchedEval=2020-04-19T14:10:59
  Partition=standard AllocNode:Sid=login2:151998
  ReqNodeList=(null) ExcNodeList=(null)
  NodeList=cn[027-028,034-035]
  BatchHost=cn027
  NumNodes=4 NumCPUs=192 NumTasks=192 CPUs/Task=1 ReqB:S:C:T=0:0:*:*
  TRES=cpu=192,node=4,billing=192
  Socks/Node=* NtasksPerN:B:S:C=48:0:*:* CoreSpec=*
  MinCPUsNode=48 MinMemoryNode=0 MinTmpDiskNode=0
  Features=(null) DelayBoot=00:00:00
  Reservation=arnab_4
  OverSubscribe=OK Contiguous=0 Licenses=(null) Network=(null)
  Command=/home/user1/slurm-job.sh
  WorkDir=/home/user1
  StdErr=/home/user1/plu.sh.e14992
  StdIn=/dev/null
  StdOut=/home/user1/plu.sh.o14992
```

Power=**scontrol update job** - change attributes of submitted job. like time limit,priority (root only)

```
$ scontrol show job 106
JobId=106 Name=slurm-job.sh
  UserId=user1(1001) GroupId=user1(1001)
  Priority=4294901717 Account=(null) QOS=normal
  JobState=RUNNING Reason=None Dependency=(null)
  Requeue=1 Restarts=0 BatchFlag=1 ExitCode=0:0
  RunTime=00:00:07 TimeLimit=14-00:00:0 TimeMin=N/A
  SubmitTime=2013-01-26T12:55:02 EligibleTime=2013-01-26T12:55:02
  StartTime=2013-01-26T12:55:02 EndTime=Unknown
  PreemptTime=None SuspendTime=None SecsPreSuspend=0
  Partition=standard AllocNode:Sid=atom-head1:3526
  ReqNodeList=(null) ExcNodeList=(null)
  NodeList=atom01
  BatchHost=atom01
  NumNodes=1 NumCPUs=2 CPUs/Task=1 ReqS:C:T=*:*:~
  MinCPUsNode=1 MinMemoryNode=0 MinTmpDiskNode=0
  Features=(null) Gres=(null) Reservation=(null)
  Shared=0 Contiguous=0 Licenses=(null) Network=(null)
  Command=/home/user1/slurm/local/slurm-job.sh
  WorkDir=/home/user1/slurm/local
```

scontrol update job= 106 TimeLimit=15-00:00:0

Suspend a job

The root user only can execute this command.

```
# scontrol suspend 135
# squeue
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
  135   standard simple.s user1 S   0:10    1    atom01
```

Resume a job

The root user only can execute this command.

```
# scontrol resume 135
# squeue
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
  135   standard simple.s user1 R   0:13    1    atom01
```

Kill a job

Users can kill their own jobs, root can kill any job.

```
$ scancel 135
$ squeue
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
```

Hold a job

```
$ squeue
JOBID PARTITION      NAME      USER  ST      TIME  NODES NODELIST(REASON)
   139      standard    simple  user1  PD       0:00      1  (Dependency)
   138      standard    simple  user1   R       0:16      1  atom01
$ scontrol hold 139
$ squeue
JOBID PARTITION      NAME      USER  ST      TIME  NODES NODELIST(REASON)
   139      standard    simple  user1  PD       0:00      1  (JobHeldUser)
   138      standard    simple  user1   R       0:32      1  atom01
```

Release a job

```
$ scontrol release 139
$ squeue
JOBID PARTITION      NAME      USER  ST      TIME  NODES NODELIST(REASON)
   139      standard    simple  user1  PD       0:00      1  (Dependency)
   138      standard    simple  user1   R       0:46      1  atom01
```

Addressing Basic Security Concerns

Your account on PARAM Brahma is 'private to you'. You are responsible for any actions emanating from your account. It is suggested that you should never share the password to anyone including your friends and system administrators!!

Please note that, by default, a new account created on PARAM Brahma is readable by everyone on the system. The following simple commands will make your account adequately safe.

<code>chmod 700 /home/\$user</code>	! will ensure that only yourself can read, write and ! execute files in your home directory
<code>chmod 750 /home/\$user</code>	! will enable yourself and the members of your ! group to read and execute files in your home ! directory
<code>chmod 755 /home/\$user</code>	! will enable yourself, your group members and ! everyone else to read and execute files in your ! directory
<code>chmod 777 /home/\$user</code>	! will enable EVERY ONE on the system to read, ! write and execute files in your home directory. ! This is a sort of 'free for all' situation. This ! should be used very judiciously

More about Batch Jobs (SLURM)

SLURM (Simple Linux Utility for Resource Management) is a workload manager that provides a framework for job queues, allocation of compute nodes, and the start and execution of jobs.

It is important to note:

- Compilations are done on the login node. Only the execution is scheduled via SLURM on the compute nodes
- Upon Submission of a Job script, each job gets a unique Job Id. This can be obtained from the '**squeue**' command.
- The Job Id is also appended to the output and error filenames.

Parameters used in SLURM job script

The job flags are used with SBATCH command. The syntax for the SLURM directive in a script is "#SBATCH <flag>". Some of the flags are used with the srun and salloc commands.

Resource	Flag Syntax	Description
partition	--partition=partition name	Partition is a queue for jobs.
time	--time=01:00:00	Time limit for the job.
nodes	--nodes=2	Number of compute nodes for the job.
cpus/cores	--ntasks-per-node=8	Corresponds to number of cores on the compute node.
resource feature	--gres=gpu:2	Request use of GPUs on compute nodes
account	--account=group-slurm-account	Users may belong to groups or accounts.
job name	--job-name="lammeps"	Name of job.
output file	--output=lammeps.out	Name of file for stdout.
	--mail-type=	Notify user by email when certain event types occur. Valid type values are NONE, BEGIN, END, FAIL, REQUEUE, ALL, TIME_LIMIT, TIME_LIMIT_90 (reached 90 percent of time limit), TIME_LIMIT_80 (reached 80 percent of time limit), and TIME_LIMIT_50 (reached 50 percent of time limit). Multiple type values may be specified in a comma separated list

email address	--mail-user=username@iiserpune.ac.in	User's email address/ User receives email notification when job states change (eg. Job completed, killed,error)
access	--exclusive	Exclusive access to compute nodes. When this access is mentioned, the complete node is dedicated to the respective job and the node is not shared with any other job

Script for a Sequential Job

```
#!/bin/bash
#SBATCH -N 1 // number of nodes
#SBATCH --ntasks-per-node=1 // number of cores per node
#SBATCH --error=job.%J.err // name of output file
#SBATCH --output=job.%J.out // name of error file
#SBATCH --time=01:00:00 // time required to execute the program
#SBATCH --partition= standard // specify the queue name
// To load the module //
module load intel/2018.5.274

cd <Path of the executable>.

/home/cdac/a.out ( Name of the executable ).
```

Script for a Parallel OpenMP Job

```
#!/bin/bash
#SBATCH -N 1 // Number of nodes
#SBATCH --ntasks-per-node=24 // Number of core per node
#SBATCH --error=job.%J.err // Name of output file
#SBATCH --output=job.%J.out // Name of error file
#SBATCH --time=01:00:00 // Time take to execute the program
#SBATCH --partition=standard // Partition or queue name

/ To load the module /
module load intel/2018.5.274 cd < path of the executable>
export OMP_NUM_THREADS=24 ( Depending upon your requirement you can change
number of threads . Maximum no.of threads is =48 )
/home/cdac/a.out (Name of the executable ).
```

Script for Parallel Job – MPI (Message Passing Interface)

```
#!/bin/sh

#SBATCH -N 16 // Number of nodes
#SBATCH --ntasks-per-node=40 // Number of cores for node
#SBATCH --time=06:50:20 // Time required to execute the program
#SBATCH --job-name=lammps // Name of application
#SBATCH --error=job.%J.err_16_node_40 // Name of the output file
```

```
#SBATCH --output=job.%J.out_16_node_40 // Name of the error file
#SBATCH --partition=standard // Partition or queue name

// To load the module //
module load intel/2018.5.274

module unload gnu8/8.3.0
source
/opt/ohpc/pub/intel2018/compilers_and_libraries_2018.5.274/linux/mkl/bin/mk
lvars.sh intel64

export I_MPI_FALLBACK=disable
export I_MPI_FABRICS=shm:dapl // Fabrics required for with node inter node
export I_MPI_DEBUG=9 // Level of debug //

cd /home/manjunath/LAMMPS/lammps-22Aug18/bench

export OMP_NUM_THREADS=1

time mpiexec.hydra -n $SLURM_NTASKS -genv OMP_NUM_THREADS 1
/home/manjunath/LAMMPS/lammps-22Aug18/src/lmp_intel_cpu_intelmpi -in in.lj
```

Script for Hybrid Parallel Job – (MPI + OpenMP)

```
#!/bin/sh
#SBATCH -N 16
#SBATCH --ntasks-per-node=40
#SBATCH --time=02:50:20
#SBATCH --job-name=lammps
#SBATCH --error=job.%J.err_16_node_40
#SBATCH --output=job.%J.out_16_node_40
#SBATCH --partition=standard

module unload gnu8/8.3.0
module load intel/2018.5.274

source
/opt/ohpc/pub/intel2018/compilers_and_libraries_2018.5.274/linux/mkl/bin/mk
lvars.sh intel64
// MPI SETTING //
export I_MPI_FALLBACK=disable
export I_MPI_FABRICS=shm:dapl // Fabrics required for with node inter node
//
export I_MPI_DEBUG=9 // Level of debug //

cd /home/manjunath/LAMMPS/lammps-22Aug18/bench

export OMP_NUM_THREADS=1 ( Depending upon your requirement you can change
number of threads . Maximum no.of threads is =48)

time mpiexec.hydra -n $SLURM_NTASKS -genv OMP_NUM_THREADS 1
/home/manjunath/LAMMPS/lammps-22Aug18/src/lmp_intel_cpu_intelmpi -in in.lj
```

I am familiar with PBS/ TORQUE. How do I migrate to SLURM?

Environment Variables	PBS/Torque	SLURM
Job Id	\$PBS_JOBID	\$SLURM_JOBID
Submit Directory	\$PBS_JOBID	\$SLURM_SUBMIT_DIR
Node List	\$PBS_NODEFILE	\$SLURM_JOB_NODELIST
Job Specification	PBS/Torque	SLURM
Script directive	#PBS	#BATCH
Job Name	-N [name]	--job-name=[name] OR -J [name]
Node Count	-1 nodes=[count]	--nodes=[min[-max]] OR -N [min[-max]]
CPU count	-1 ppn=[count]	---ntasks-per-node=[count]
CPUs Per Task		--cpus-per-task=[count]
Memory Size	-1 mem-[MB]	--mem=[MB] OR – mem_per_cpu=[MB]
Wall Clock Limit	-1 walltime=[hh:mm:ss]	--time=[min] OR – mem_per_cpu=[MB]
Node Properties	-1 nodes=4.ppn=8:[property]	--constraint=[list]
Standard Output File	-o [file_name]	--output=[file_name] OR -o [file_name]
Standard Error File	-e [file_name]	--error=[file_name] OR -e {file_name}
Combine stdout/stderr	-j oe (both to stdout)	(This is default if you do not specify –error)
Job Arrays	-t [array_spec]	--array=[array_spec] OR -a [array_spec]
Delay Job Start	-a [time]	--begin=[time]

Preparing your own Executable

The compilations are done on the login node, whereas the execution happens on the compute nodes via the scheduler (SLURM).

Note: The Compilation and execution must be done with same libraries and matching version to avoid unexpected results.

Steps:

1. Load required modules on the login node.
2. Do the compilation.
3. Open the job submission script and specify the same modules to be loaded as used while compilation.
4. Submit the script.

The directory contains a few sample programs and their sample job submission scripts. The compilation and execution instructions are described in the beginning of the respective files.

The user can copy the directory to his/her home directory and further try compiling and executing these sample codes. The command for copying is as follows:

```
cp -r /home/apps/cdac/samples/ ~/.
```

1. mm.c - Serial Version of Matrix-Matrix Multiplication of two NxN matrices
2. mm_omp.c - Basic OpenMP Version of Matrix-Matrix Multiplication of two NxN matrices
3. mm_mpi.c - Basic MPI Version of Matrix-Matrix Multiplication of two NxN matrices
4. mm_acc.c - OpenAcc Version of Matrix-Matrix Multiplication of two NxN matrices
5. mm_blas.cu - CUDA Matrix Multiplication program using the CuBlas library.
6. mm_mkl.c - MKL Matrix Multiplication program.
7. laplace_acc.c - OpenACC version of the basic stencil problem.

It is recommended to use the intel compilers since they are better optimized for the hardware.

Compilers

Compilers	Description	Versions Available
gcc/gfortran	GNU Compiler (C/Fortran)	4.8.5, 8.3.0
icc/ifort	Intel Compiler (C/Fortran)	17.x, 18.x
mpicc/mpif90	Intel-mpi based GNU compiler (C/Fortran)	
mpiicc/mpiifort	Intel-mpi based intel compiler (C/Fortran)	
nvcc	CUDA C Compiler	7.0,8.0,9.0,9.2,10.0,10.1
pgcc/pgf90	PGI Compiler (C/Fortran)	19.4, 2019

Optimization Flags

Optimization flags are meant for uniprocessor optimization, wherein, the compiler tries to optimize the program, on the basis of the level of optimization. The optimization flags can be explored more on the respective compiler pages. A few examples are given below.

```
Intel: -O3 -xHost
GNU: -O3
PGI: -fast
```

Given next is a brief description of compilation and execution of the various types of programs. However, for certain bigger applications, loading of additional dependency libraries might be required.

C Program:

```
Setting up of environment: module load intel/2018.5.274
compilation: icc -O3 -xHost <<prog_name.c>>
Execution: ./a.out
```

C+OpenMP Program:

```
Setting up of environment: module load intel/2018.5.274 compilation: icc -
O3 -xHost -qopenmp <<prog_name.c>>
Execution: ./a.out
```

C+MPI Program:

```
Setting up of environment: module load intel/2018.5.274 compilation: mpiicc
-O3 -xHost <<prog_name.c>>
Execution: mpirun -n <<num_procs>> ./a.out
```

C+MKL Program:

```
Setting up of environment: module load intel/2018.5.274 compilation: icc -
O3 -xHost -mkl <<prog_name.c>>
Execution: ./a.out
```

CUDA:

```
Setting up of environment: module load cuda/10.1
compilation: nvcc <<prog_name.cu>> -lcublas -ar=sm_70
Execution: ./a.out
```

OpenACC:

```
Setting up of environment: module load pgi/2019/pgi/19.4
Compilation: pgcc -acc -fast -Minfo=all -ta=tesla:managed -Mprof=ccff
<<prog_name.c>>
Execution: ./a.out
```

Job Submission on Scheduler (SLURM):

A sample job submission scripts for each of the sample programs is given. Upon completion/termination of the execution, two files (output and error) are generated.

A few sample commands for SLURM are as follows:

sinfo	Lists out the status of resources in the system
squeue	Lists out the Job information in the system
sbatch <<job_script>>	Submitting a job to the scheduler

scancel <<job_name>>	Delete a job
-------------------------	--------------

Job Scheduling on PARAM Brahma

Scheduler

PARAM Brahma has Slurm-19.05.0-1 (open source) as a workload manager for HPC facility. Slurm is a widely used batch scheduler in top500 HPC list. PARAM Brahma consists of two types of compute nodes: i.e. CPU only (192 GB) nodes, High memory (768 GB) nodes.

Following partitions/queues have been defined for different requirements

1. **standard**: CPU, High memory Jobs.
2. **cpu**: cpu only jobs
3. **hm**: High memory intensive jobs

All users can submit to the Standard partition. The standard Partition contains CPU, high memory nodes. If user wants to submit a job only on cpu nodes, he/she can use cpu partition. If user wants to submit a job only on high memory, he/she can use hm partition.

Resource limits like priority etc., will be defined as per client requirements.

sinfo

This Slurm command is used to view available **partition** and **node** information on the cluster.

```
[root@brahma ~]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
standard*  up 4-00:00:00    3  down* cn[009,102,189]
standard*  up 4-00:00:00    3   mix  cn[084-086]
standard*  up 4-00:00:00   150  alloc cn[001-002,014-017,022-023,025,027-028,036-083,087-101,103-172,186,190-194]
standard*  up 4-00:00:00    23  resv  cn[003-008,010-013,018-021,024,026,029-035]
standard*  up 4-00:00:00    3   idle cn[185,187-188]
cpu        up 4-00:00:00    2  down* cn[102,189]
cpu        up 4-00:00:00    3   mix  cn[084-086]
cpu        up 4-00:00:00   135  alloc cn[040-083,087-101,103-172,186,190-194]
cpu        up 4-00:00:00    3   idle cn[185,187-188]
hm         up 4-00:00:00    1  down* cn009
hm         up 4-00:00:00   15  alloc cn[001-002,014-017,022-023,025,027-028,036-039]
hm         up 4-00:00:00    23  resv  cn[003-008,010-013,018-021,024,026,029-035]
[root@brahma ~]#
```

Figure 9 – View of available partition and node information on the cluster

Walltime

Walltime parameter defines as to how long your job will run. The maximum runtime of a job allowed is 4 days. If more than 4 days are required, a special request needs to be sent to HPC coordinator and it will be dealt with on a case to case basis. The command line to specify walltime is given below.

```
srun -t walltime <days-hours:mins:seconds>
```

and also as part of the submit scripts described in the manual. If a job does not get completed within the walltime specified in the script, it will get terminated.

The biggest advantage of specifying appropriate walltime is that the efficiency of scheduling improves resulting in improved throughput in all jobs including yours. You are encourage to arrive at the appropriate walltime for your job by executing your jobs few times.

NOTE: You are requested to explicitly specify the walltime in your command lines and scripts.

Per user

- Every user will have quota of 500 GB in HOME file system (/home) and 1TB in SCRATCH file system
- Users are recommended to copy their execution environment and input files to scratch file system (/scratch/<username>) during job running and copy output data back to HOME area
- Maximum number of running jobs at a time per account is 10
- Each account has been given GRPTRESMINs=32*24*60 Mins (46080) mins

Per Department

Each department has been allocated equal fare share. All IISER Pune department are listed below:

- bio
- chemistry
- ecs
- iiser
- math
- physics

Scheduling Type

PARAM Brahma has been configured with Slurm's backfill scheduling policy. It is good for ensuring higher system utilization; it will start lower priority jobs if doing so does not delay the expected start time of any higher priority jobs. Since the expected start time of pending jobs depends upon the expected completion time of running jobs, reasonably accurate time limits are important for backfill scheduling to work well.

JOB PRIORITY

The job's priority at any given time will be a weighted sum of all the factors that have been enabled in the slurm.conf file. Job priority can be expressed as:

```
Job_priority =
    (PriorityWeightAge) * (age_factor) +
    (PriorityWeightFairshare) * (fair-share_factor) +
    (PriorityWeightJobSize) * (job_size_factor) +
    (PriorityWeightPartition) * (partition_factor) +
    (PriorityWeightQOS) * (QOS_factor) +
    SUM(TRES_weight_cpu * TRES_factor_cpu,
        TRES_weight_<type> * TRES_factor_<type>,
        ...)
```

All of the factors in this formula are floating point numbers that range from 0.0 to 1.0. The weights are unsigned, 32 bit integers. The job's priority is an integer that ranges between 0 and 4294967295. The larger the number, the higher the job will be positioned in the queue, and the sooner the job will be scheduled. A job's priority, and hence its order in the queue, can vary over time. For example, the longer a job sits in the queue, the higher its priority will grow when the age weight is non-zero.

Age Factor: The age factor represents the length of time a job has been sitting in the queue and eligible to run. Current value for Age factor is 10000.

Job Size Factor: The job size factor correlates to the number of nodes or CPUs the job has requested. Current value for Job Size factor is 1000.

Partition Factor: Each node partition can be assigned an integer priority. The larger the number, the greater the job priority will be for jobs that request to run in this partition. Current value for partition factor is 15000.

Quality of Service (QOS) Factor: Each QOS can be assigned an integer priority. The larger the number, the greater the job priority will be for jobs that request this QOS. Current value for QOS factor is 1000.

Fair-share Factor: The fair-share component to a job's priority influences the order in which a user's queued jobs are scheduled to run based on the portion of the computing resources

they have been allocated and the resources their jobs have already consumed. Current value for fair-share factor is 100000.

SSHARE

This tool is for listing the shares of association to a cluster.

```
[root@brahma ~]# sshare
```

Account	User	RawShares	NormShares	RawUsage	EffectvUsage	FairShare
root			0.000000	8030415018	1.000000	
root	root	1	0.009901	714907	0.000089	1.000000
iiser		60	0.594059	7589633739	0.945111	
bio		1	0.166667	199917169	0.026341	
collins		1	0.200000	199651766	0.998672	
madhu		1	0.200000	264922	0.001325	
sanjeev		1	0.200000	480	0.000002	
suhita		1	0.200000	0	0.000000	
chemistry		1	0.166667	3047567778	0.401543	
aloke		1	0.333333	308447386	0.101211	
arnab		1	0.333333	1562903032	0.512836	
arun		1	0.333333	1176217358	0.385953	
ecs		1	0.166667	780622152	0.102854	
neena		1	0.500000	1775	0.000002	
shyam		1	0.500000	780620376	0.999998	
math		1	0.166667	0	0.000000	
anindya		1	1.000000	0	0.000000	
physics		1	0.166667	3560805299	0.469167	
mukul		1	0.200000	455868399	0.128024	
prasenjit		1	0.200000	3079405537	0.864806	
santhanam		1	0.200000	485949	0.000136	
sreejith		1	0.200000	0	0.000000	
nsmapp		20	0.198020	440066372	0.054800	
cdac		1	0.250000	31893353	0.072474	
others		1	0.066667	16770908	0.525843	
ews		1	0.250000	408147284	0.927468	
nsmddp		1	0.250000	25734	0.000058	
nsmext		20	0.198020	0	0.000000	

Figure 10 - Listing the shares of association to a cluster

ACCOUNTING

Accounting system tracks and manages HPC resource usage. As jobs are completed or resources are utilized, accounts are charged and resource usage is recorded. Accounting policy is like a bank/Credit System, where each department can be allocated with some pre-defined budget on a quarterly basis for CPU usage. As and when the resources are utilized, the amount will be deducted. The allocation will be reset at end of every quarter.

sacct

This command can report resource usage for running or terminated jobs including individual tasks, which can be useful to detect load imbalance between the tasks.

sstat

This command can be used to status only currently running jobs. It also can give you

sreport

This command can be used to generate reports based upon all jobs executed in a particular time interval.

Debugging Your Codes

Introduction

A **debugger** or **debugging tool** is a computer program that is used to test and debug other programs (the "target" program).

When the program "traps" or reaches a preset condition, the debugger typically shows the location in the original code if it is a source-level debugger or symbolic debugger, commonly now seen in **integrated development environments**.

Debuggers also offer more sophisticated functions such as running a program step by step (single-stepping or program animation), stopping (breaking) (pausing the program to examine the current state) at some event or specified instruction by means of a breakpoint, and tracking the values of variables.

Some debuggers have the ability to modify program state while it is running. It may also be possible to continue execution at a different location in the program to bypass a crash or logical error.

Basics How Tos

Compilation

Compilation with a separate flag '-g' is required since the program needs to be linked with debugging symbols.

```
gcc -g <program_name.c>  
e.x. gcc -g random_generator.c
```

Running with gdb:

gdb is a command line utility available with almost all Linux systems' compiler collection packages.

```
gdb <executable.out>  
e.x. gdb a.out
```

Basic gdb commands (to be executed in gdb command line window):**Start:**

Starts the program execution and stops at the first line of the main procedure. Command line arguments may be provided if any.

Run:

Starts the program execution but does not stop. It stops only when any error or program trap occurs. Command line arguments may be provided if any.

Help:

Prints the list of command available. Specifying 'help' followed by a command (e.x. 'help run') displays more information about that command.

File <filename>:

Loads a binary program that is compiled with '-g' flag for debugging.

List [line_no]

Displays the source code (nearby 10 lines) of the program in execution where the execution stopped. If 'line_no' is specified, it display the source code (10 lines) at the specified line.

Info:

Displays more information about the set of utilities and saved information by the debugger. For example; 'info breakpoints' will list all the breakpoints, similarly 'info watchpoints' will list all the watchpoints set by the user while debugging their programs.

Print <expression>:

Prints the values of variables / expression at the current running instance of the program.

Step N:

Steps the program one (or 'N') instructions ahead or till the program stops for any reason. Steps through each and every instruction even if it is function call (only function or instruction compiled with debugging flags).

next:

This command also steps through the instructions of the program. Unlike 'step' command, if the current source code line calls a subroutine, this command does not enter the subroutine, but instead steps over the call, in effect treating it as a single source line.

Continue:

This command continues the stopped program till the next breakpoint has occurred or till the end of the program. It is used to continue from a paused/debug point state.

Break [sourcefile:]<line_no> [if condition]:

Stops the program at the specified line number and provides a breakpoint for the user. Specific source code file and breakpoint based on a condition can also be set for specific cases. You can also view the list of breakpoints set, by using the 'info breakpoints' command.

watch <expression>:

A watchpoint means break the program or stop the execution of the program when the value of the expression provided is changed. Using watch command specific variables can be watched for value changes. You can also view the list of watchpoints by using the 'info watchpoints' command.

Delete <breakpoint number>

Delete command deletes a breakpoint or a watchpoint that has been set by a user while debugging the program.

Backtrace:

Prints the backtrace of all stack frames of the program. Provides the call stack and more other information about the running program.

These are some of the most powerful utilities that can be used to debug your programs using gdb. gdb is not limited to these commands and contains a rich set of features that can allow you to debug multi-threaded programs as well. Also, all the commands, along with the ones listed above have 'n' number of different variants for more in-depth control. Same can be utilized using the help page of gdb.

Using gdb (example – inspecting the code)

For this case study, we have a small program that generates a long unique random number for each run.

Let's look at the code we have.

```

#include <stdio.h>           //printf
#include <stdlib.h>          //malloc, srand, rand
#include <unistd.h>          //getpid

#define N 100
#define N_LEN 100

//Generate a short random number
short rand_fract(void) {
    short sum = 1;
    for (short i = 0; i < (rand() % N); ++i)
        for (short j = 0; i < N; ++j) {
            int value = (i * j) / (i + j);
            sum += (value != 0) ? value : sum;
        }
    return sum;
}

//Returns the factorial of a number
long long factorial(unsigned int x) {
    if (x == 1 || x == 0)
        return 1LL;
    else
        return (x * factorial(x - 1));
}

```

Figure 11 – snapshot of debugging process

Things to note:

- 1) We have a few libraries included for the functions that are used in the program.
- 2) We have two ‘#define’ statements:
 - a. ‘N’ for the number of times the ‘rand_fract’ function will spend in calculating the random number.
 - b. ‘N_LEN’ for the length of the final random number string generated. Currently it is set to ‘100’ which means that the long random number will be of length 100.
- 3) Then, we have a function by name ‘rand_fract’ that iterates over two loops and using the values of iterators (‘i’ and ‘j’), it calculates a small random number. Since, ‘rand()’ function is used for the outer loop, its number of iterations cannot be clearly defined which gives the function a random nature.
- 4) The next function is as simple as its name is. It just takes an unsigned integer and returns its factorial.

PART 2:

```

int main (int argc, char *argv[]) {

    short f1 = 0;

    //Create a random seed based on process id.
    srand((unsigned int) getpid());

    //Generate a random number salt.
    f1 = rand_fract() % 10;

    //Get the factorial of the number
    long long random_fact = factorial(f1);

    //Normalize the factorial to number modulo N_LEN + 1
    int normalized_fact = random_fact % (N_LEN + 1);

    int *array = NULL;

    //Create an array of size obtained from normalized factorial modulo N_LEN + 1
    array = (int *) malloc (sizeof (int) * normalized_fact);
    if (array == NULL) { printf("Not enough memory\n"); return -1; }

    //Populate the array with integers in reverse order
    //Double the number five times if it is even
    for (int i = 0; i < normalized_fact; ++i) {
        array[i] = (normalized_fact - i);
    }

    //Print the serial number
    for (int i = normalized_fact - 1; i >= 0; --i)
        printf("%0d", (array[i] + rand()) % 10);
    for (int i = (N_LEN - normalized_fact); i > 0; --i)
        printf("%0d", (rand() % 10));
    printf("\n");

    //Free allocated memory
    free(array);

    return 0;
}

```

Figure 12 – Snapshot of debugging process

Things to note:

- 1) This is the main function of the program.
- 2) The flow of the main function is as follows:
 - a. The program first sets a random seed using the process-id of the program.
 - b. It calls 'rand_fract' function and the resultant random number is operated by a modulo 10 operation. Finally, the result is stored in the variable 'f1'.
 - c. Next the factorial of the obtained 'f1' is calculated and stored in 'random_fact'.
 - d. This result is again passed through a modulo 'N_LEN + 1' and stored in 'normalized_fact'.
 - e. Then a dynamic array is constructed and partially filled with integer values in descending order from the 'normalized_fact' value.

- f. Finally, the partial array is printed by mixing the value of the array with rand() function values followed by a modulo 10 operation.
- g. The remaining partial part of final random value is generated using a basic rand() modulo 10 operation.

Using gdb (example – using the debugger)

The code that we looked upon seems correct, as well as it compiles successfully without any errors. But, when we run this code snippet, this is the result we get.

```
$ gcc random_generator.c
$ ./a.out
Floating point exception (core dumped)
$ █
```

Figure 13- output at a debugging stage

The program ended up with a core dump without giving much information but just 'Floating point exception'. Now let's compile the code with debugging information and run the program simply with gdb.

```
$ gcc -g random_generator.c
$ gdb a.out
GNU gdb (GDB) Fedora 8.3.50.20190824-25.fc31
Copyright (C) 2019 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
  <http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from a.out...
(gdb) set style enabled off
(gdb) run
Starting program: /home/vineetm/debugger/a.out

Program received signal SIGFPE, Arithmetic exception.
0x00000000004011cc in rand_fract () at random_generator.c:13
13          int value = (i * j) / (i + j);
(gdb) █
```

Figure 14 – Snapshot of debugging process

Here we compiled the code using '-g' and then used the 'run' command we studied earlier for running the program. You can observe that the debugger stopped at line number 13 where the 'Floating point exception (SIGFPE)' occurred. At this point we can even go and check the code at line number 13. But for now, let's check what other information we can get from the debugger. Let's check the values of the variables 'i' and 'j' at this point.

```
(gdb) run
Starting program: /home/vineetm/debugger/a.out

Program received signal SIGFPE, Arithmetic exception.
0x00000000004011cc in rand_fract () at random_generator.c:13
13                               int value = (i * j) / (i + j);
(gdb) print i
$1 = 0
(gdb) print j
$2 = 0
(gdb) █
```

Figure 15 – Output depicting “Arithmetic Exception”

The values of both 'i' and 'j' appear to be '0' and thus a **divide by zero** exception is what caused our program to terminate. Let's update the code such that the value of 'i' and 'j' will never become '0'. This is the modified code:

```
//Generate a short random number
short rand_fract(void) {
    short sum = 1;
    for (short i = 1; i < (rand() % N); ++i)
        for (short j = 1; j < N; ++j) {
            int value = (i * j) / (i + j);
            sum += (value != 0) ? value : sum;
        }
    return sum;
}
```

Figure 16 – Snapshot of debugging process

Thus, we just updated the loop index variables to start from '1' instead of '0'. **Thus, using gdb, it was very simple to identify the point where the error occurred.** Let's re-run our updated code and check what we get.

```
$ gcc random_generator.c
$ ./a.out
Floating point exception (core dumped)
$ █
```

Figure 17 – Well, we dumped core !!

WHAT!? This is unexpected. We just cured the error part of our program and still getting an FPE. Let's go through the debugger and check where the error point is right now.

```
$ gcc -g random_generator.c
$ gdb a.out
GNU gdb (GDB) Fedora 8.3.50.20190824-25.fc31
Copyright (C) 2019 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
    <http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from a.out...
(gdb) set style enabled off
(gdb) run
Starting program: /home/vineetm/debugger/a.out

Program received signal SIGFPE, Arithmetic exception.
0x00000000004011cc in rand_fract () at random_generator.c:13
13      int value = (i * j) / (i + j);
(gdb) print i
$1 = 1
(gdb) print j
$2 = -1
(gdb) list
8      //Generate a short random number
9      short rand_fract(void) {
10         short sum = 1;
11         for (short i = 1; i < (rand() % N); ++i)
12             for (short j = 1; i < N; ++j) {
13                 int value = (i * j) / (i + j);
14                 sum += (value != 0) ? value : sum;
15             }
16         return sum;
17     }
(gdb) █
```

Figure 18 - Snapshot of debugging process

The debugger output shows that the error occurred on the same line as earlier. But in this case, the value of 'i' and 'j' are not '0,0' but they are '1, -1' which is causing the denominator at line 13 to be '0' and thus, causing an FPE. In addition to print commands, we have also issued the 'list' command which shows the nearby 10 lines of the code where the program stopped.

You can observe that some bugs in the programs are easier to debug but some aren't.

We will have to dig in much more to find out what is going on. Also, to be noted, we have our inner loop iterating from 1 to N (which is 100), but still the value of 'j' is printed out to be '-1'. How is this even possible!? Smart programmers would have the problem identified, but let's stick to the basics on how to gdb. Let us use the 'break' command and set a breakpoint at line number 13 and observe what is going on.

```
(gdb) list 13
8      //Generate a short random number
9      short rand_fract(void) {
10         short sum = 1;
11         for (short i = 1; i < (rand() % N); ++i)
12             for (short j = 1; i < N; ++j) {
13                 int value = (i * j) / (i + j);
14                 sum += (value != 0) ? value : sum;
15             }
16         return sum;
17     }
(gdb) break 13
Breakpoint 1 at 0x4011b5: file random_generator.c, line 13.
(gdb) info breakpoints
Num      Type          Disp Enb Address              What
1        breakpoint     keep y   0x00000000004011b5 in rand_fract at random_generator.c:13
(gdb) run
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/vineetm/debugger/a.out

Breakpoint 1, rand_fract () at random_generator.c:13
13                 int value = (i * j) / (i + j);
(gdb) print i
$3 = 1
(gdb) print j
$4 = 1
(gdb) █
```

Figure 19 – setting breakpoint

Thus, using the command 'break 13' we have set the breakpoint at line number 13 which was verified using the 'info breakpoint' command. Then, we reran the program with the 'run' command. At line 13 the program stopped and using 'print' command we checked the values of 'i' and 'j'. At this point, all seems to be well. Now, let's proceed further. For stepping 1 instruction we can use the 'step' command. Let's do that and observe the value of 'j'.

```

(gdb) print j
$5 = 1
(gdb) step
14             sum += (value != 0) ? value : sum;
(gdb) step
12             for (short j = 1; i < N; ++j) {
(gdb) step

Breakpoint 1, rand_fract () at random_generator.c:13
13             int value = (i * j) / (i + j);
(gdb) print j
$6 = 2
(gdb) step
14             sum += (value != 0) ? value : sum;
(gdb) step
12             for (short j = 1; i < N; ++j) {
(gdb) step

Breakpoint 1, rand_fract () at random_generator.c:13
13             int value = (i * j) / (i + j);
(gdb) print j
$7 = 3
(gdb) █

```

Figure 20 – single stepping through to catch error !!

You can observe the usage of the 'step' command. We are going through the program line by line and checking the values of the variable 'j'.

There seems to be a lot of writing/typing of the 'step' command just to proceed with the program. Since, we have already set a breakpoint at line 13, we can use another command called as 'continue'. This command continues the program till the next breakpoint or the end of the program.

```

(gdb) continue
Continuing.

Breakpoint 1, rand_fract () at random_generator.c:13
13          int value = (i * j) / (i + j);
(gdb) print j
$8 = 4
(gdb) continue
Continuing.

Breakpoint 1, rand_fract () at random_generator.c:13
13          int value = (i * j) / (i + j);
(gdb) print j
$9 = 5
(gdb) continue
Continuing.

Breakpoint 1, rand_fract () at random_generator.c:13
13          int value = (i * j) / (i + j);
(gdb) print j
$10 = 6
(gdb) █

```

Figure 21 – Debugging continued

You can see that we reduced the typing of 'step' command by 3 times to a 'continue' command just 1 time. But this is also having us write 'continue' and 'print' multiple times. Let us use some other utility in gdb known as 'data breakpoints' also known as watchpoints. But before that, let us delete the existing breakpoint using the 'delete' command.

```

(gdb) info breakpoints
Num      Type             Disp Enb Address                  What
1        breakpoint      keep y   0x0000000000004011b5 in rand_fract at random_generator.c:13
        breakpoint already hit 6 times
(gdb) delete 1
(gdb) info breakpoints
No breakpoints or watchpoints.
(gdb) █

```

Figure 22 – Debugging continued

Now let us see how to set a watchpoint.

```
(gdb) watch j
Hardware watchpoint 2: j
(gdb) info watchpoints
Num      Type          Disp Enb Address      What
2        hw watchpoint keep y                j
(gdb) continue
Continuing.

Hardware watchpoint 2: j

Old value = 6
New value = 7
0x00000000004011f5 in rand_fract () at random_generator.c:12
12                for (short j = 1; i < N; ++j) {
(gdb)
Continuing.

Hardware watchpoint 2: j

Old value = 7
New value = 8
0x00000000004011f5 in rand_fract () at random_generator.c:12
12                for (short j = 1; i < N; ++j) {
(gdb)
Continuing.

Hardware watchpoint 2: j

Old value = 8
New value = 9
0x00000000004011f5 in rand_fract () at random_generator.c:12
12                for (short j = 1; i < N; ++j) {
(gdb) █
```

Figure 23 – Setting a watch point

Thus, using the command 'watch j' we have set a watchpoint over 'j'. Now every time when the value of 'j' changes, a break will occur. You can also note the old and new values of 'j' printed out at each break. Another point to note is that after having one 'continue' command, the program had a break. Further, by just pressing the 'Enter/Return' button on the keyboard, the continue command was repeated. Thus, by pressing the 'Enter/Return' button, the last command is repeated. At this point, we have learned much about the debugger, but we are still not able to proceed fast with our error. Is there any other way to proceed? Well, yes!!

We want to observe at the point where the value of 'j' reaches closer to 'N i.e. 100'. Which means that we are only concerned about what happens after 'j' reaches 99. Here, we land up on using what is called as conditional breakpoints. First, we will delete our watchpoint and then make use of the conditional breakpoint.

```
(gdb) info watchpoints
Num      Type          Disp Enb Address      What
2        hw watchpoint keep y          j
        breakpoint already hit 4 times
(gdb) delete 2
(gdb) list 13
8          //Generate a short random number
9          short rand_fract(void) {
10             short sum = 1;
11             for (short i = 1; i < (rand() % N); ++i)
12                 for (short j = 1; i < N; ++j) {
13                     int value = (i * j) / (i + j);
14                     sum += (value != 0) ? value : sum;
15                 }
16             return sum;
17         }
(gdb) break random_generator.c:13 if j == 99
Note: breakpoint 3 also set at pc 0x4011b5.
Breakpoint 4 at 0x4011b5: file random_generator.c, line 13.
(gdb) continue
Continuing.

Breakpoint 3, rand_fract () at random_generator.c:13
13                     int value = (i * j) / (i + j);
(gdb) print j
$12 = 99
(gdb) █
```

Figure 24 – Debugging continued

You can observe another variant of the 'break' command. We have explicitly stated the file and the line number along with a condition to stop. This is useful, when the source code is large and having multiple files. After setting a conditional break, we stopped at the point where the value of 'j' becomes '99'. Now, let us see what happens next. Since, this is a critical point at which we could observe the program, it is better if we step in the program using the 'step' command instead of relying on any break/watch points.

```

(gdb) print j
$17 = 99
(gdb) step
14                sum += (value != 0) ? value : sum;
(gdb)
12                for (short j = 1; i < N; ++j) {
(gdb)
13                int value = (i * j) / (i + j);
(gdb) print j
$18 = 100
(gdb) step
14                sum += (value != 0) ? value : sum;
(gdb)
12                for (short j = 1; i < N; ++j) {
(gdb)
13                int value = (i * j) / (i + j);
(gdb) print j
$19 = 101
(gdb) █

```

Figure 25 – Well, Back to square one !!

This, is unexpected!! The value of 'j' should never be 100 or anything above it.

Thus, something is wrong with the conditional statement!!

By observation, we have figured out that the condition is itself wrong. It should have been ' $j < N$ ' instead of ' $i < N$ '. This is a silly mistake of the programmer that lead us to this much of an effort.

Also, the value of 'j' which was observed as '-1' was an outcome of the 'short' datatype overflow i.e. the value of 'j' went from 1 to 32767 (assuming short as 2 bytes) and then from -32768 to -1.

Finally, a hard programming bug was discovered. Let us correct this error and rerun the program.

```

$ gcc random_generator.c
$ ./a.out
Segmentation fault (core dumped)
$ ./a.out
1648815196934936907712847411075269363872465178968652936899126642679968327854843818024803725602089977
$ ./a.out
Segmentation fault (core dumped)
$ ./a.out
5697819555377639608368302418588269943918647330492449391532502328545856833586737093122407957112268963
$ ./a.out
6150930494475890863050318719122734582864309765193799040843958123681888230308039318234438024068348747
$ ./a.out
Segmentation fault (core dumped)
$ █

```

Figure 26 – Again Dumping Core!! Things are getting interesting or frustrating or both!!

This is strange!!

Sometimes the program is getting the correct output, but sometimes, we are getting a segmentation fault. Debugging such a program may be tricky since the occurrence of the bug is low. We will proceed with our standard debugger steps to identify the error.

```
$ gcc -g random_generator.c
$ gdb a.out
GNU gdb (GDB) Fedora 8.3.50.20190824-25.fc31
Copyright (C) 2019 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
  <http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from a.out...
(gdb) set style enabled off
(gdb) run
Starting program: /home/vineetm/debugger/a.out
5411371059776263605409873043180521681086975694174815924540859823191291008689026600122878853935366497
[Inferior 1 (process 61832) exited normally]
(gdb)
```

Figure 27 – Debugging continued

We compiled the code and ran it using the debugger. But the program completed successfully. Let us rerun it till a point where the program fails.

```
(gdb) run
Starting program: /home/vineetm/debugger/a.out
5411371059776263605409873043180521681086975694174815924540859823191291008689026600122878853935366497
[Inferior 1 (process 61832) exited normally]
(gdb) run
Starting program: /home/vineetm/debugger/a.out
7919846386128432134671007571802513619267000845358948048917009272836772572766214308134147179016591178
[Inferior 1 (process 61978) exited normally]
(gdb) run
Starting program: /home/vineetm/debugger/a.out

Program received signal SIGSEGV, Segmentation fault.
0x000000000040126c in factorial (x=4294792703) at random_generator.c:24
24         return (x * factorial(x - 1));
(gdb)
```

Figure 28 – Debugging continued

Here we observe a point where the program exited at the function ‘factorial’.

This is a point where the debugger didn’t give much information about what the value of the variable ‘x’ was. It just pointed out that the program failed at the function named ‘factorial’. That’s it!

Another reason for such kind of output would be because of the recursive nature of the function. The stack frame where the function ‘factorial’ failed could be in a long nest of recursive calls. At such points, it would be better to inspect the program at an earlier point

and look for errors. Let us have a breakpoint before the 'factorial' function was called and view the value of the parameters that are passed to the function.

```
(gdb) list main
22         return 1LL;
23     else
24         return (x * factorial(x - 1));
25     }
26
27     int main (int argc, char *argv[]) {
28
29         short f1 = 0;
30
31         //Create a random seed based on process id.
(gdb)
32         srand((unsigned int) getpid());
33
34         //Generate a random number salt.
35         f1 = rand_fract() % 10;
36
37         //Get the factorial of the number
38         long long random_fact = factorial(f1);
39
40         //Normalize the factorial to number modulo N_LEN + 1
41         int normalized_fact = random_fact % (N_LEN + 1);
(gdb) break 36
Breakpoint 1 at 0x4012da: file random_generator.c, line 38.
(gdb) run
Starting program: /home/vineetm/debugger/a.out

Breakpoint 1, main (argc=1, argv=0x7fffffff0d8) at random_generator.c:38
38         long long random_fact = factorial(f1);
(gdb) print f1
$1 = 1
(gdb) continue
Continuing.
9962554943440906583333593426000827274699155147995250801174774876796185292736525250533642241728519329
[Inferior 1 (process 62328) exited normally]
(gdb) █
```

Figure 29 – Debugging continued (Will it ever end?)

Thus, we have set a breakpoint before the call of the function 'factorial' and ran the program. For the value of 'f1 = 8' for the 'factorial' function the process seems to exit normally. Let us rerun.

```
(gdb) run
Starting program: /home/vineetm/debugger/a.out

Breakpoint 1, main (argc=1, argv=0x7fffffff0d8) at random_generator.c:38
38         long long random_fact = factorial(f1);
(gdb) print f1
$1 = -8
(gdb) continue
Continuing.

Program received signal SIGSEGV, Segmentation fault.
0x000000000040126c in factorial (x=4294792699) at random_generator.c:24
24         return (x * factorial(x - 1));
(gdb) █
```

Figure 30 – We are almost there !!

Unexpectedly, we have got the value of 'f1' as '-8' and the program seems to have crashed. Let us observe the 'rand_fract' function and 'factorial' function once again. And study the behavior of the functions where we could get a negative number.

```
(gdb) list rand_fract
4
5     #define N 100
6     #define N_LEN 100
7
8     //Generate a short random number
9     short rand_fract(void) {
10         short sum = 1;
11         for (short i = 1; i < (rand() % N); ++i)
12             for (short j = 1; j < N; ++j) {
13                 int value = (i * j) / (i + j);
(gdb)
14                 sum += (value != 0) ? value : sum;
15             }
16         return sum;
17     }
18
19     //Returns the factorial of a number
20     long long factorial(unsigned int x) {
21         if (x == 1 || x == 0)
22             return 1LL;
23         else
(gdb) run
Starting program: /home/vineetm/debugger/a.out

Breakpoint 1, main (argc=1, argv=0x7fffffffed0d8) at random_generator.c:38
38         long long random_fact = factorial(f1);
(gdb) print f1
$2 = -8
(gdb) █
```

Figure 31 – Debugging continued

Important points here to observe are:

The 'rand_fract' function is returning a datatype of 'short' while the calculation of the return value could be significantly large which may overflow the size of 'short', thus, causing a negative answer.

The function 'factorial' is expecting a value of type 'unsigned int'. Since the value passed to the function is a negative value, having an implicit conversion from a negative number to an unsigned number means that we are having a very large value passed to the factorial function.

Also, since the 'factorial' function is recursive, passing a very large number to it could cause multiple calls to the same function and thus, overflowing the stack provided to the user.

Now let us, step further into our program and see whether what we are discussing is the same behavior that is being observed.

```
(gdb) print f1
$4 = -8
(gdb) step
factorial (x=4294967288) at random_generator.c:21
21         if (x == 1 || x == 0)
(gdb)
24         return (x * factorial(x - 1));
(gdb)
factorial (x=4294967287) at random_generator.c:21
21         if (x == 1 || x == 0)
(gdb)
24         return (x * factorial(x - 1));
(gdb)
factorial (x=4294967286) at random_generator.c:21
21         if (x == 1 || x == 0)
(gdb)
24         return (x * factorial(x - 1));
(gdb)
factorial (x=4294967285) at random_generator.c:21
21         if (x == 1 || x == 0)
(gdb)
24         return (x * factorial(x - 1));
(gdb)
factorial (x=4294967284) at random_generator.c:21
21         if (x == 1 || x == 0)
(gdb) █
```

Figure 32 – At last a clue!!!

This is what we had expected!!

A number '-1' passed to the 'factorial' function is being implicitly converted to a very large number '4294967295'.

Stepping in more reveals the recursive behavior of the 'factorial' function i.e. each call is having a sub call to the same function with one value less. Thus, what to do in these types of cases. Assume you have a large code where these functions are called from multiple locations.

Modifying the signature of any of the function means changing the code everywhere where the function is called. This is not affordable!! These are some cases, where a choice is to be made where patching the code is necessary for semantics of the program.

Let us observe a piece of code where this change can be made and then test our program for the expected results.

```
int main (int argc, char *argv[]) {  
  
    short f1 = 0;  
  
    //Create a random seed based on process id.  
    srand((unsigned int) getpid());  
  
    //Generate a random number salt.  
    f1 = rand_fract() % 10;  
  
    f1 = abs(f1);  
  
    //Get the factorial of the number  
    long long random_fact = factorial(f1);  
  
    //Normalize the factorial to number modulo N_LEN + 1  
    int normalized_fact = random_fact % (N_LEN + 1);  
  
    int *array = NULL;
```

Figure 33 - Correction applied !!

By observing the code, we find out that the expected value of 'f1' is between '0 to 9' (because of the modulo 10 operation).

Thus, without changing the signature of any function, we have inserted a patch (the highlighted) portion, that maintains the semantics of the code as well cures the problem that we had. Now let us just run and check our final program.

```

$ gcc random_generator.c
$ ./a.out
1947155904444356260827867895829013940560127574392384362061544757042318542200659899527743928595211645
$ ./a.out
0929989745546167856100961512939018573760223504833542534886091294243732854126729096261941760801537820
$ ./a.out
0244202592758390536991444038465396583053516022410228562188134665524049393105566500577005828487059653
$ ./a.out
2872718293228567054539368096969066437379893671576029177909795701346393295764931536773483363035181911
$ ./a.out
9128766061538956027759598074797832715087451437704122190965898083361413690723150214543517739636518290
$ ./a.out
4700580792312412551673394453147630608790931492649027378923259025287077290331618510470262819931652479
$ ./a.out
3597977632870365479023130705918446909083470263354375991983675631252252710058384841530848408963208645
$ ./a.out
0864510419056291282368845079139095210792697191764209304803037158672651132052448868790301906812889064
$ ./a.out
4972916609538445900529958158240849030612776510222275380497441425328380877450674923651890544608240290
$ ./a.out
9528608642866177753983842182285047120984190000785095691019238964676666205506776407087180325311790389
$ █

```

Figure 34 – Resolved !!!

Thus, we are getting the correct results as expected.

Conclusions

We started with a program that we assumed to be functional but then the program ended up with bugs that were not straightforward. We then explored the power of the debugger and the various ways to identify the bugs in our program. We looked upon the easy solutions, and slowly migrated towards the type of bugs that are not easily traceable.

Finally, we identified and corrected all the bugs in our program with the help of the debugger and arrived at a bug free code.

Points to Note

- Bugs in the program cannot be necessarily a compilation error.
- One type of error can be caused by multiple bugs in the same line of code.
- Sometimes, it is not possible to change the code even when the problem is identified. The best way to cure this is to study the behavior of the code and apply patches wherever necessary.
- Using simple utilities from the ‘GNU Debugger’ can help in getting rid of problem causing bugs in large programs.

Overall Coding Modifications Done

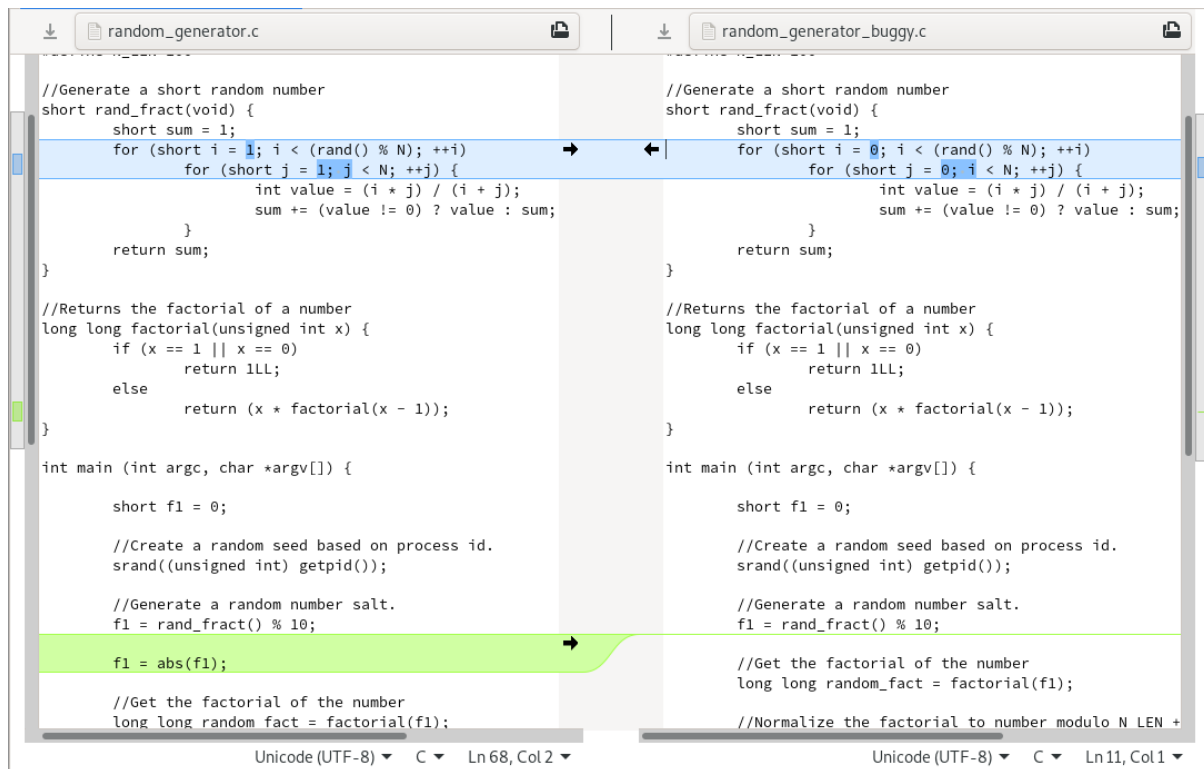


Figure 35 – What all we did to get things right !

Machine Learning / Deep Learning Application Development

Most of the popular python based machine learning/deep learning libraries are installed on PARAM Brahma system. While developing and testing their applications, users have option to choose different environment / runtime setup like “virtual environment-based python libraries” or “conda runtime based python libraries”.

For most of the major environment (virenv, conda) different modules are prepared. Users can check the list of the modules by using “**module avail**” command. Shown below is an example of loading conda environment in current bash shell and continue with application development.

Once logged into PARAM Brahma HPC Cluster, check which all libraries are available, loaded in current shell. To check list of modules loaded in current shell, use the command given below:

```
$ module list
```

To check all modules available on the system, but not loaded currently, use the command given below:

```
$ module avail
```

To activate conda environment on PARAM Brahma, load module “conda-python/3.7” as shown below:

```
$ module load conda-python/3.7
```

Conda environment has been installed with most of the popular python packages as shown below

Tensorflow	Tensorflow-gpu	Mpi4py	Keras
Theano	Scipy	Scikit-Learn	Pytorch

Once “conda-python/3.7” module is loaded, end-users can use all libraries inside their python program. Many other modules based on virtual env are available on the system.

Users can load those libraries using “module load” command and use them for their applications.

How to Install your own Software?

There are two approaches to install software.

1. System wide installation
2. Local installation.

System wide installation can be done by only admin. If you wish to do this, please approach system administrator. User can do local installation in their home directory. In this section we are describing the installation of HMMER application in user’s home directory.

Local installation

Step 1. Login to Brahma cluster by using your credential.

Step 2. Download the software that you want to install . For example to download HMMER software use the command given below.

```
$ wget http://eddylib.org/software/hmmer/hmmer.tar.gz
```

Step 3. Untar the file. (if your software in zip format use unzip command)

```
$ tar xzf hmmer.tar.gz
```

Step 4. go to the software folder.

```
$ cd hmmer-3.3
```

Step 5. configure the installation path.

```
$ ./configure --prefix /your/install/path
```

Step 6. now run the 'make' command for install the software on installation path.

```
$ make
```

The newly compiled binaries are now in the src directory.

Step 7. Runs a test suite that checks for errors in the software (optional)

```
$ make check
```

Step 8. run 'make install' to install the programs and man pages in your location mention in step 2 #

```
$ make install
```

By default, programs are installed in **/usr/local/bin** and man pages in **/usr/local/share/man/man1/**, if you do not provide installation path in step 2.

* This is general instruction for installation, please refer the installation instruction or manual or readme file that comes with software for more details.

if you get any dependency error, resolve that or ask system admin to install that dependency if not installed.

Reference link: <http://hmmer.org/documentation.html>

Some Important Facts

About File Size

The global/home is served by a number of storage arrays. Each of the storage array contains a portion of the global/home. The size of a disk in the storage array is 2TB (2000 GB).

Technically, the size of a file can be about 2000 GB (which is really big). However, since the disk is shared by a large number of files, effectively the size of a single file will be far smaller.

Normally, this file size is kept to be about few GBs which is sufficient for most of the users.

However, if you wish to have file sizes which are larger than this, you need to create files ACROSS disks and this process is known as ‘striping’.

```
lfs setstripe -c 4 .
```

After this has been done all new files created in the current directory will be spread over 4 storage arrays each having 1/4th of the file. The file can be accessed as normal no special action needs to be taken. When the striping is set this way, it will be defined on a per directory basis so different directories can have different stripe setups in the same file system, new subdirectories will inherit the striping from its parent at the time of creation.

We recommend users to set the stripe count so that each chunk will be approx. 200-300GB each, for example

File Size	Stripe count	Command
500-1000 GB	4	<code>lfs setstripe -c 4 .</code>
1000 – 2000 GB	8	<code>lfs setstripe -c 8</code>

Once a file is created with a stripe count, it cannot be changed. A user by themselves are also able to set stripe size and stripe count for their directories and A user can check the set stripe size and stripe count with command:

```
lfs getstripe <path to the direcorey>
```

To set the stripe count as

```
lfs setstripe -c 4 -s 10m <path to the direcorey>
```

The options on the above command used have these respective functions.

- **-c** to set the stripe count; 0 means use the system default (usually 1) and -1 means stripe over all available OSTs (lustre Object Storage Targets).
- **-s** to set the stripe size; 0 means use the system default (usually 1 MB) otherwise use k, m or g for KB, MB or GB respectively

Little-Endian and Big-Endian issues?

By and large, most of the computers follow little-endian format. This essentially means that the last byte of the binary representation of data is stored first. However, there is another way of representing data (used in some machines) where in the first byte of the binary representation of data is stored first. When binary files are to be read across these different kinds of machines, bytes need to be re-ordered. Many compilers do support this feature. Please explore this aspect, if a perfectly working code on a given machine, fails to get executed on another machine (with a different processor).

Best Practices for HPC

1. Do **NOT** run any job which is longer than a few minutes on the login nodes. Login node is for compilation of job. It is best to run the job on compute nodes. (compute nodes)
2. It is **recommended** to go through the beginner's guide in **/home/apps/cdac/samples**. This should serve as a good starting point for the new users.
3. Use the same compiler to compile different parts/modules/library-dependencies of an application. Using different compilers (e.g. pgcc + icc) to compile different parts of application may cause linking or execution issues.
4. Choosing appropriate compiler switches/flags/options (e.g. -O3) may increase the performance of application substantially (accuracy of output must be verified). Please refer to documentation of compilers (online / docs present inside compiler installation path / man pages etc.)
5. Modules/libraries used for execution should be the same as that used for compilations. This can be specified in the Job submission script.
6. Be aware of the amount of disk space utilized by your job(s). Do an estimate before submitting multiple jobs.
7. Please submit jobs preferably in \$SCRATCH. You can back up your results/summaries in your \$HOME
8. \$SCRATCH is NOT backed up! Please download all your data to your Desktop/ Laptop.
9. Before installing any software in your home, ensure that it is from a reliable and safe source. Ransomware is on the rise!
10. Please do not use spaces while creating the directories and files.
11. Please inform PARAM Brahma support when you notice something strange - e.g. unexpected slowdowns, files missing/corrupted etc.

Installed Applications/Libraries

Following is the list of few of the applications from various domains of science and engineering installed in the system.

HPC Applications	Bio-informatics	MUMmer, HMMER, MEME, Schrodinger, PHYLIP, mpiBLAST, ClustalW,
	Molecular Dynamics	NAMD (for CPU and GPU), LAMMPS, GROMACS
	Material Modeling, Quantum Chemistry	Quantum-Espresso, Abinit, CP2K, NWChem,
	CFD	OpenFOAM, SU2
	Weather, Ocean, Climate	WRF-ARW, WPS (WRF), ARWPost (WRF), RegCM, MOM, ROMS
Deep Learning Libraries	cuDNN, TensorFlow, Tensorflow with Intel Python , Tensorflow with GPU, Theano, Caffe , Keras , numpy, Scipy, Scikit-Learn, pytorch.	
Visualization Programs	GrADS, ParaView, VisIt, VMD	
Dependency Libraries	NetCDF, PNETCDF, Jasper, HDF5, Tcl, Boost, FFTW	

Standard Application Programs on PARAM Brahma

The purpose of this section is to expose the users to different application packages which have been installed. Users interested in exploring these packages may kindly go through the scripts, typical input files and typical output files. It is suggested that, at first, the users may submit the scripts provided and get a feel of executing the codes. Later, they may change the parameters and the script to meet their application requirements.

LAMMPS Applications

LAMMPS is an acronym for **L**arge-scale **A**tomic/ **M**olecular **M**assively **P**arallel **S**imulator. This is extensively used in the fields of Material Science, Physics, Chemistry and may others.

More information about LAMMPS may please be found at <https://lammps.sandia.gov> .

1. The LAMMPS input is **in.lj** file which contains the below parameters.

Input file = in.lj

```
# 3d Lennard-Jones melt

variable      x index 1
variable      y index 1
variable      z index 1

variable      xx equal 64*$x
variable      yy equal 64*$y
variable      zz equal 64*$z

units         lj
atom_style    atomic

lattice       fcc 0.8442
region        box block 0 ${xx} 0 ${yy} 0 ${zz}
create_box    1 box
create_atoms   1 box
mass          1 1.0

velocity      all create 1.44 87287 loop geom

pair_style     lj/cut 2.5
pair_coeff     1 1 1.0 1.0 2.5

neighbor       0.3 bin
neigh_modify   delay 0 every 20 check no

fix           1 all nve

run           1000000
```

2. THE LAMMPS RUNNING SCRIPT

```
#!/bin/sh

#SBATCH -N 32
#SBATCH --ntasks-per-node=40
#SBATCH --time=02:50:20
#SBATCH --job-name=lammps
#SBATCH --error=job.%J.err_32_node_40
#SBATCH --output=job.%J.out_32_node_40
#SBATCH --partition=standard
```

```

module unload gnu8/8.3.0
module load intel/2018.5.274

source
/opt/ohpc/pub/intel2018/compilers_and_libraries_2018.5.274/linux/mkl/bin/mk
lvars.sh intel64

export I_MPI_FALLBACK=disable
export I_MPI_FABRICS=shm:dapl
export I_MPI_DEBUG=9

cd /home/manjunath/LAMMPS/lammps-22Aug18/bench

export OMP_NUM_THREADS=1

time mpiexec.hydra -n $SLURM_NTASKS -genv OMP_NUM_THREADS 1
/home/manjunath/LAMMPS/lammps-22Aug18/src/lmp_intel_cpu_intelmpi -in in.lj

```

3. LAMMPS OUTPUT FILE.

```

LAMMPS (22 Aug 2018)
  using 1 OpenMP thread(s) per MPI task
Lattice spacing in x,y,z = 1.6796 1.6796 1.6796
Created orthogonal box = (0 0 0) to (107.494 107.494 107.494)
  8 by 10 by 16 MPI processor grid
Created 1048576 atoms
  Time spent = 0.048476 secs
Neighbor list info ...
  update every 20 steps, delay 0 steps, check no
  max neighbors/atom: 2000, page size: 100000
  master list distance cutoff = 2.8
  ghost atom cutoff = 2.8
  binsize = 1.4, bins = 77 77 77
  1 neighbor lists, perpetual/occasional/extra = 1 0 0
  (1) pair lj/cut, perpetual
    attributes: half, newton on
    pair build: half/bin/atomonly/newton
    stencil: half/bin/3d/newton
    bin: standard
Setting up Verlet run ...
  Unit style      : lj
  Current step    : 0
  Time step       : 0.005
Per MPI rank memory allocation (min/avg/max) = 2.699 | 2.703 | 2.708 Mbytes
Step Temp E_pair E_mol TotEng Press
    0          1.44   -6.7733681          0   -4.6133701   -5.0196704
 1000000    0.65695755  -5.7125359          0   -4.7271005    0.48799127
Loop time of 723.716 on 1280 procs for 1000000 steps with 1048576 atoms

Performance: 596918.946 tau/day, 1381.757 timesteps/s
99.5% CPU use with 1280 MPI tasks x 1 OpenMP threads

MPI task timing breakdown:
Section | min time | avg time | max time | %varavg | %total
-----

```

Pair	424.38	435.47	461.05	26.2	60.17
Neigh	59.782	60.365	62.991	3.9	8.34
Comm	193.24	219.39	231.11	38.5	30.31
Output	0.00013494	0.00085223	0.0088639	0.0	0.00
Modify	6.4813	6.6462	7.541	5.6	0.92
Other		1.841			0.25

Nlocal: 819.2 ave 845 max 786 min
 Histogram: 3 2 34 115 256 372 315 137 33 13
 Ngghost: 2417.97 ave 2468 max 2369 min
 Histogram: 8 31 81 216 314 327 202 76 22 3
 Neighs: 30698 ave 32432 max 28796 min
 Histogram: 4 16 47 194 306 325 245 103 34 6

Total # of neighbors = 39293494
 Ave neighs/atom = 37.4732
 Neighbor list builds = 50000
 Dangerous builds not checked
 Total wall time: 0:12:25

OpenFOAM

openFOAM is a free, opensource software which covers most areas of Engineering and Science. It can be used to solve very interesting problems in fields ranging from Turbulence, Heat transfer, Acoustics, Electromagnetics, complex fluid flows including chemical reactions, solid mechanics and a lot more. Please follow the link <https://www.openfoam.com> to get more information.

Description of Inputs for openFOAM

Input file is taken from NASA website which does wing body simulation. The data can be copied from /home/apps/Data/OpenFOAM path on PARAM Brahma

Grid size: 10 million

Solver: sonicFoam

Iterations: 4000

Decomposition of grid is done using Metis.

Script of OpenFOAM

```
#!/bin/sh
#SBATCH -N 14
#SBATCH --ntasks-per-node=48
#SBATCH --threads-per-core=1
#SBATCH --ntasks=672
#SBATCH --time=06:50:20
#SBATCH --job-name=openfoam
#SBATCH --error=job.%J.err_16_node_40
#SBATCH --output=job.%J.out_16_node_40
#SBATCH --partition=cpu

ulimit -s unlimited
#module load intel/2018.0.1.163

module unload gnu8/8.3.0

export I_MPI_FALLBACK=disable
export I_MPI_FABRICS=shm:dapl
export I_MPI_DEBUG=5
export I_MPI_PIN_PROCESSOR_LIST=0-39
#####MXM Optimization#####
export I_MPI_DAPL_SCALABLE_PROGRESS=1
export I_MPI_RDMA_TRANSLATION_CACHE=1
export I_MPI_FAIR_CONN_SPIN_COUNT=2147483647
export I_MPI_FAIR_READ_SPIN_COUNT=2147483647
#export I_MPI_ADJUST_REDUCE 2, I_MPI_ADJUST_BCAST 0
export I_MPI_RDMA_TRANSLATION_CACHE=1
export I_MPI_RDMA_RNDV_BUF_ALIGN=65536
export I_MPI_SPIN_COUNT=121
```

```

export I_MPI_DAPL_DIRECT_COPY_THRESHOLD=65536
#export I_MPI_DAPL_UD=enable
#####
source /home/shwetad/OpenFOAM/Intel-2018/openfoam_bashrc_2018

export OMP_NUM_THREADS=1
cd /home/shwetad/OpenFOAM_DATA/NSM

rm -rf processor*
decomposePar
(time mpirun -np 672 sonicFoam -parallel) 2>&1 | tee out_4000_NSM_14Node

Output Values after 4000 iterations:
forceCoeffs forces write:
    Cm      = -8.50123
    Cd      = 0.0327941
    Cl      = -1.926
    Cl(f)   = -9.46423
    Cl(r)   = 7.53823

```

The said iterations complete in 7minutes 55 seconds on 14 nodes.

WRF Application

The **Weather Research and Forecasting (WRF)** Model is a next-generation mesoscale numerical weather prediction system designed to serve both operational forecasting and atmospheric research needs. WRF is suitable for a broad spectrum of applications across scales ranging from meters to thousands of kilometers. WRF was developed at the National Center for Atmospheric Research (NCAR) which is operated by the University Corporation for Atmospheric Research (UCAR), USA.

More information about WRF may please be found at:

<https://www.mmm.ucar.edu/weather-research-and-forecasting-model>

For a reference run, the dataset used is as following with model simulation time being reduced to 15 minutes:

Dataset: Single domain, large size. 2.5 km CONUS, June 4, 2005

(Ref: https://www2.mmm.ucar.edu/wrf/WG2/benchv3/#_Toc212961289)

The WRF input files used for reference run are present in */home/apps/Data/WRF/input/run*

```
# Changes/Suggestions to namelist.input

&time_control
  run_hours           = 0,
  run_minutes         = 15,
  io_form_history     = 11,           //parallel-netcdf
  io_form_restart     = 11,
  io_form_input       = 11,
  io_form_boundary    = 11,           //serial-netcdf
  io_form_auxhist2    = 2,

&dynamics
  use_baseparam_fr_nml = .t.,

&namelist_quilt           // For no. of nodes (e.g. greater than 32
                           nodes) using quilt servers gives better
                           performance
  nio_tasks_per_group = 0,
  nio_groups = 1,
```

1. WRF job submission SLURM script

The following reference job script is placed in
/home/apps/Data/WRF/input/run/wrf_4n.sh

```
#!/bin/bash
```

```

#SBATCH -N 4
#SBATCH --ntasks-per-node=40
#SBATCH --time=00:30:00
#SBATCH --job-name=WRF_CONUS
#SBATCH --error=job.%J.err
#SBATCH --output=job.%J.out
#SBATCH --partition=cpu

cd $SLURM_SUBMIT_DIR

###Loading WRF environment
module load wrf/3.8.1/intel2018

###Creating list of nodes to map WRF MPI processes
mpirun.hydra -n $SLURM_NTASKS hostname > hosts.txt
sort -u hosts.txt > hosts_wrf.txt
sed -i 's/$/:20/' hosts_wrf.txt

###Two OpenMP threads per MPI rank
WRFMPI=`expr $SLURM_NTASKS / 2 `

###Setting Intel MPI environment
export I_MPI_DEBUG=9
export I_MPI_FALLBACK=disable

      (time mpirun.hydra --machinefile hosts_wrf.txt -env I_MPI_PIN_DOMAIN
omp:compact -env OMP_NUM_THREADS=2 -env KMP_STACKSIZE=200m -n $WRFMPI
wrf.exe) >& 4n.2omp.wrf.out

### To save execution command to out file
echo "(time mpirun.hydra --machinefile hosts_wrf.txt -env I_MPI_PIN_DOMAIN
omp:compact -env OMP_NUM_THREADS=2 -env KMP_STACKSIZE=200m -n $WRFMPI
wrf.exe)" >> 4n.2omp.wrf.out

```

2. WRF Output Snippet

```

$tail rsl.out.0000
Timing for main: time 2005-06-04_06:14:45 on domain 1: 2.16898 elapsed seconds
Timing for main: time 2005-06-04_06:15:00 on domain 1: 2.16480 elapsed seconds
wrf: SUCCESS COMPLETE WRF

```

The above workload on four compute nodes took approximate execution time of 3min15sec

NAMD Application

Nano Scale Molecular Dynamics (NAMD) software for molecular dynamics simulation is designed for high-performance simulations of Large Macro Molecular system on parallel computers. This software also makes use of GPGPUs.

More information about NAMD may please be found at

<http://www.ks.uiuc.edu/Research/namd/>

Citation:

"NAMD was developed by the Theoretical and Computational Biophysics Group in the Beckman Institute for Advanced Science and Technology at the University of Illinois at Urbana-Champaign."

James C. Phillips, Rosemary Braun, Wei Wang, James Gumbart, Emad Tajkhorshid, Elizabeth Villa, Christophe Chipot, Robert D. Skeel, Laxmikant Kale, and Klaus Schulten.
Scalable molecular dynamics with NAMD. *Journal of Computational Chemistry*, 26:1781-1802, 2005.

For a reference run, NAMD version 2.13 with following dataset has been used:

Dataset: STMV (virus) benchmark (1,066,628 atoms, periodic, PME)

(Ref: <https://www.ks.uiuc.edu/Research/namd/utilities/>)

Please follow the examples to get familiarized with writing scripts for using NAMD with CPU cores and GPGPUs.

The NAMD input files and SLURM job scripts used for reference runs are present in */home/apps/reference/namd/*

1. Contents of *stmv.namd* file used for a reference run with *numsteps* increased to 5000

```
#####
## ADJUSTABLE PARAMETERS ##
#####
structure      stmv.psf
coordinates     stmv.pdb

#####
## SIMULATION PARAMETERS ##
#####

# Input
paraTypeCharmm      on
parameters          par_all127_prot_na.inp
temperature         298

# Force-Field Parameters
exclude             scaled1-4
1-4scaling          1.0
cutoff              12.
switching           on
switchdist          10.
pairlistdist        13.5

# Integrator Parameters
timestep            1.0
nonbondedFreq       1
fullElectFrequency  4
```



```

stepspercycle      20

# Constant Temperature Control
langevin           on      ;# do langevin dynamics
langevinDamping    5       ;# damping coefficient (gamma) of 5/ps
langevinTemp       298
langevinHydrogen   off     ;# don't couple langevin bath to hydrogens

# Constant Pressure Control (variable volume)
useGroupPressure   yes ;# needed for rigidBonds
useFlexibleCell    no
useConstantArea    no

langevinPiston     on
langevinPistonTarget 1.01325 ;# in bar -> 1 atm
langevinPistonPeriod 100.
langevinPistonDecay 50.
langevinPistonTemp 298

cellBasisVector1   216.832  0.  0.
cellBasisVector2   0.  216.832  0.
cellBasisVector3   0.  0  216.832
cellOrigin         0.  0.  0.

PME                on
PMEGridSizeX       216
PMEGridSizeY       216
PMEGridSizeZ       216

# Output
outputName         stmv-output

outputEnergies     20
outputTiming       20

numsteps           5000

```

2. NAMD job submission SLURM script for GPGPUs

The reference job script is placed at `/home/apps/reference/namd/namd_gpu_2.13.sh`

```

#!/bin/bash

#SBATCH -N 5
#SBATCH --ntasks-per-node=40
#SBATCH --time=00:10:00
#SBATCH --job-name=NAMD_GPU
#SBATCH --error=job.%J.err
#SBATCH --output=job.%J.out
#SBATCH --partition=gpu

cd $SLURM_SUBMIT_DIR
module load namd/2.13/mpi2019v5/cuda

```

```

module unload gnu8/8.3.0
module load intel/2019.5.281

mpiexec.hydra -n $SLURM_NTASKS hostname > hosts.txt
sort -u hosts.txt > hosts_namd.txt
sed -i 's/$/:2/' hosts_namd.txt

NAMDMPI=`expr $SLURM_NTASKS / 20 `
export I_MPI_DEBUG=9
#export I_MPI_FABRICS=shm:dapl
export FI_PROVIDER=mlx
export I_MPI_FALLBACK=0

      (time mpiexec.hydra --machinefile hosts_namd.txt -n $NAMDMPI namd2
+ppn 19 +pemap 1-19,21-39 +commap 0,20 +setcpuaffinity +isomalloc_sync
+idlepoll ./stmv.namd +devices 0,1) >& 5N.namd.cuda.out
echo "(time mpiexec.hydra --machinefile hosts_namd.txt -n $NAMDMPI namd2
+ppn 19 +pemap 1-19,21-39 +commap 0,20 +setcpuaffinity +isomalloc_sync
+idlepoll ./stmv.namd +devices 0,1)" >> 5N.namd.cuda.out
mpiicc -v          >> 5N.namd.cuda.out
icc -v            >> 5N.namd.cuda.out
which namd2       >> 5N.namd.cuda.out

```

3. NAMD job submission SLURM script for CPUs

The reference job script is placed at */home/apps/reference/namd/namd_cpu_2.13.sh*

```

#!/bin/bash

#SBATCH -N 16
#SBATCH --ntasks-per-node=40
#SBATCH --time=00:15:00
#SBATCH --job-name=NAMD
#SBATCH --error=job.%J.err
#SBATCH --output=job.%J.out
#SBATCH --partition=cpu

cd $SLURM_SUBMIT_DIR
module load namd/2.13/mpi2019v5/cpu
module unload gnu8/8.3.0
module load intel/2019.5.281

mpiexec.hydra -n $SLURM_NTASKS hostname > hosts.txt
sort -u hosts.txt > hosts_namd.txt
sed -i 's/$/:2/' hosts_namd.txt

NAMDMPI=`expr $SLURM_NTASKS / 20 `
export I_MPI_DEBUG=9
#export I_MPI_FABRICS=shm:dapl
export FI_PROVIDER=mlx
export I_MPI_FALLBACK=0
mpiicc -v

      (time mpiexec.hydra --machinefile hosts_namd.txt -n $NAMDMPI namd2
+ppn 19 +pemap 1-19,21-39 +commap 0,20 +setcpuaffinity +isomalloc_sync
+idlepoll ./stmv.namd) >& 16N.namd.cpu.out

```

```
echo "(time mpiexec.hydra --machinefile hosts_namd.txt -n $NAMDMPID namd2
+ppn 19 +pemap 1-19,21-39 +commap 0,20 +setcpuaffinity +isomalloc_sync
+idlepoll ./stmv.namd)" >> 16N.namd.cpu.out
mpiicc -v >> 16N.namd.cpu.out
icc -v >> 16N.namd.cpu.out
which namd2 >> 16N.namd.cpu.out
```

4. NAMD Output Snippet

```
TIMING: 5000 CPU: 53.5804, 0.010235/step Wall: 53.6822, 0.0102851/step, 0 hours remaining,
5024.433594 MB of memory in use.
ETITLE: TS BOND ANGLE DIHED IMPRP
ELECT VDW BOUNDARY MISC KINETIC TOTAL
TEMP POTENTIAL TOTAL3 TEMPAVG PRESSURE GPRESSURE
VOLUME PRESSAVG GPRESSAVG
ENERGY: 5000 368697.5464 279967.9454 81941.7526 5087.5336 -
4524058.5568 385721.1155 0.0000 0.0000 945098.3696 -2457544.2937
297.2577 -3402642.6633 -2449021.7220 297.4163 -83.0695 -13.9811
10199588.8034 13.0980 11.6157
```

5. The above workload took approximate execution times as:

S. No.	Resource Type	No. of Compute Nodes	ns/day	Execution Time (sec)
1	GPGPU (with CPU)	5	8	75
2	CPU	16	8	75

GROMACS Application

GROningen MACHine for Chemical Simulations (GROMACS) is a [molecular dynamics](#) package mainly designed for simulations of [proteins](#), [lipids](#), and [nucleic acids](#). It was originally developed in the Biophysical Chemistry department of [University of Groningen](#), and is now maintained by contributors in universities and research centres worldwide. GROMACS is one of the fastest and most popular software packages available, and can run on [central processing units](#) (CPUs) and [graphics processing units](#) (GPUs).

Input description of Gromacs

Input file can be download from

ftp://ftp.gromacs.org/pub/benchmarks/water_GMX50_bare.tar.gz

The mdp option used is pme with 50000 steps

```
#!/bin/sh
#SBATCH -N 10
```

```
#SBATCH --ntasks-per-node=48
##SBATCH --time=03:05:30
#SBATCH --job-name=gromacs
#SBATCH --error=job.16.%J.err
#SBATCH --output=job.16.%J.out
#SBATCH --partition=cpu

cd /home/shweta/Gromacs/water-cut1.0_GMX50_bare/3072
source /home/shweta/gromcs_openmpi_3_2_0/bin/GMXRC
export I_MPI_DEBUG=5
ulimit -s unlimited
export OMP_NUM_THREADS=1
export PATH=/home/apps/openmpi_3.2.0_gcc_4.8.5/bin:$PATH
export
LD_LIBRARY_PATH=/home/apps/openmpi_3.2.0_gcc_4.8.5/lib:/home/apps/openmpi_3
.2.0_gcc_4.8.5/lib64:$LD_LIBRARY_PATH
export
INCLUDE_PATH=/home/apps/openmpi_3.2.0_gcc_4.8.5/include:$INCLUDE_PATH
mpirun -np 4 gmx_mpi grompp -f pme.mdp -c conf.gro -p topol.top
(time mpirun -np $SLURM_NTASKS gmx_mpi mdrun -s topol.tpr) 2>&1 | tee
log_gromacs_40_50k_mpirun
```

Output:

```
Number of logical cores detected (40) does not match the number reported by
OpenMP (1).
Consider setting the launch configuration manually!
Running on 10 nodes with total 480 cores, 480 logical cores
Cores per node:          48
Logical cores per node:  48
Hardware detected on host cn177 (the node of MPI rank 0):
CPU info:
Vendor: GenuineIntel
Brand:  Intel(R) Xeon(R) Gold 6148 CPU @ 2.40GHz
SIMD instructions most likely to fit this hardware: AVX2_256
SIMD instructions selected at GROMACS compile time: AVX2_256
Reading file /home/shwetad/Gromacs/water-cut1.0_GMX50_bare/3072/topol.tpr,
VERSION 5.1.4 (single precision)
Changing nstlist from 10 to 20, rlist from 1 to 1.032
Will use 392 particle-particle and 128 PME only ranks
This is a guess, check the performance at the end of the log file
Using 480 MPI processes
Using 1 OpenMP thread per MPI process
Non-default thread affinity set probably by the OpenMP library,
disabling internal thread affinity
Back Off! I just backed up ener.edr to ./#ener.edr.9#
starting mdrun 'Water'
50000 steps,    100.0 ps.
Writing final coordinates.
Back Off! I just backed up confout.gro to ./#confout.gro.7#
Average load imbalance: 4.7 %
Part of the total run time spent waiting due to load imbalance: 3.4 %
Average PME mesh/force load: 1.018
Part of the total run time spent waiting due to PP/PME imbalance: 1.1 %
Core t (s)    Wall t (s)          (%)
```

```
Time:      245555.569      472.286      51993.0
          (ns/day)      (hour/ns)
Performance:      18.294      1.312
The above input takes 7minutes 55 seconds on 10 nodes.
```

Acknowledging the National Supercomputing Mission in Publications

If you use supercomputers and services provided under the National Supercomputing Mission, Government of India, please let us know of any published results including Student Thesis, Conference Papers, Journal Papers and patents obtained.

Please acknowledge the National Supercomputing Mission as given below:

The support and the resources provided by PARAM Brahma Facility under the National Supercomputing Mission, Government of India at the Indian Institute of Science Education and Research; Pune are gratefully acknowledged.

Also, please submit the copies of dissertations, reports, reprints and URLs in which “National Supercomputing Mission, Government of India” is acknowledged to:

HoD HPC Technologies,
Centre for Development of Advanced Computing,
CDAC Innovation Park,
S.N. 34/B/1,
Panchavati, Pashan,
Pune – 411008
Maharashtra

Communication of your achievements using resources provided by National Supercomputing Mission, will help the Mission in measuring outcomes and gauging the future requirements. This will also help in further augmentation of resources at a given site of National Supercomputing Mission.

Getting Help – PARAM Brahma Support

We suggest that you kindly visit the link <https://parambrahma.iiserpune.ac.in/FAQ> to know if the problem you are facing has already been reported by some other user and has already been solved. If this is the case, you will find a summary of the solution provided by the system administrator there. If you do not find any mention of the issue being faced by you, please refer to these four easy steps to generate a Ticket related to the issue you are experiencing.

Your Ticket will be assisted by the Brahma Support team. The ticket generated will be closed only when the related issue gets resolved.

You can generate a new ticket for any of the new issue that you are experiencing.

Steps to Create a New Ticket

1. Place the URL (<http://parambrahma.iiserpune.ac.in/support/>) in your browser.
2. On the right-top corner of the page click **Sign In**. Refer to Fig: 1 for the same.

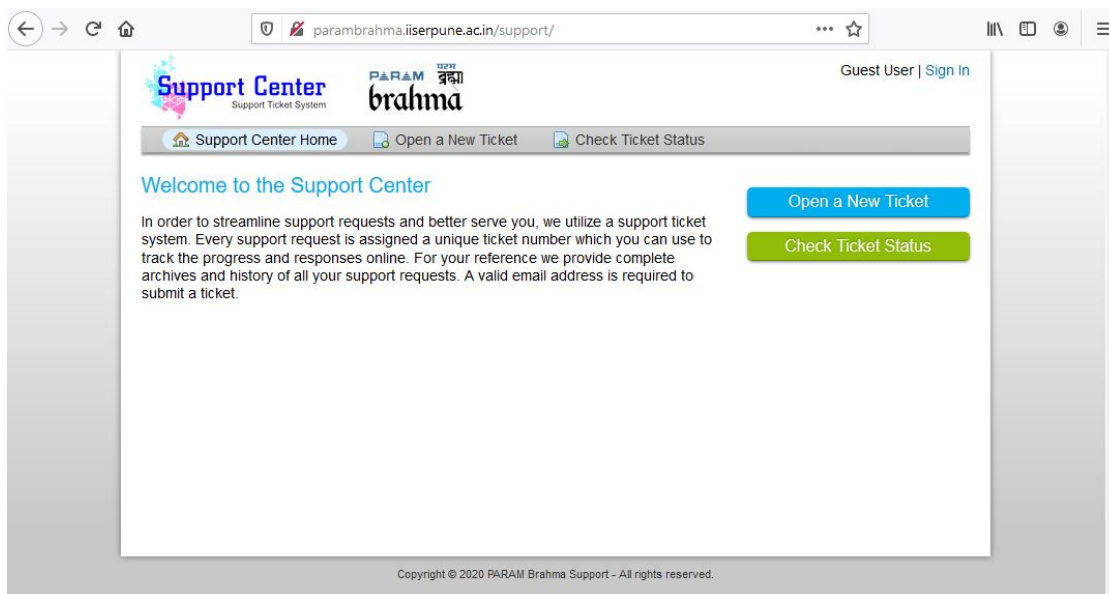


Figure 36 – Snapshot of Ticketing System

3. Sign in by using the Username and Password that you use for logging to the Cluster.
Refer to Fig2 for the same.

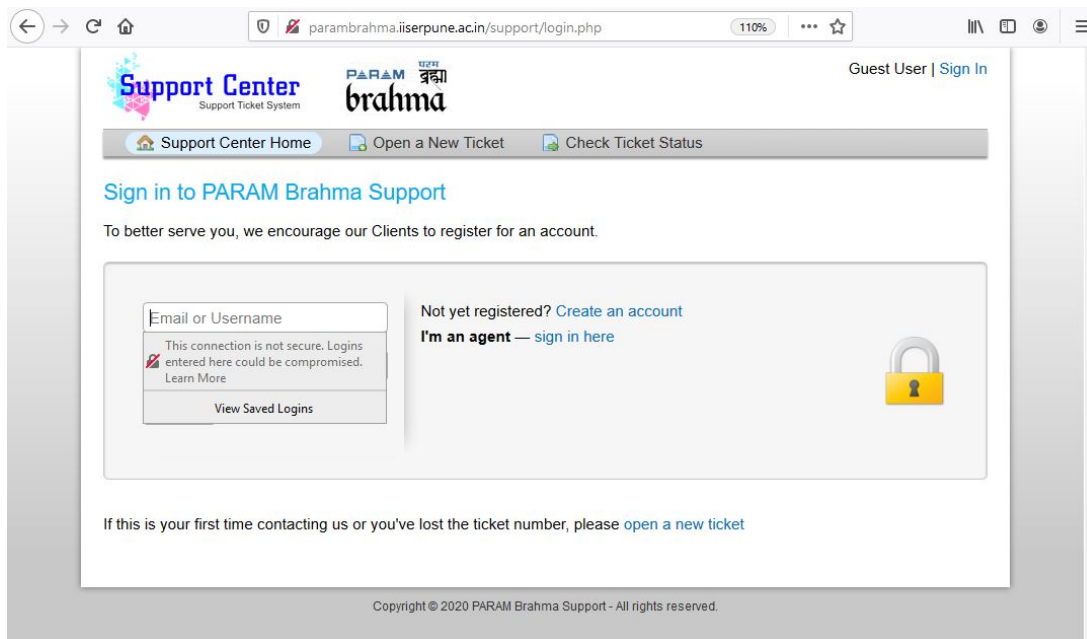


Figure 37 - Snapshot of Ticketing System

4. Select a **Help Topic** from the Dropdown and then Click on **Create Ticket**. Refer to Fig:3 for the same

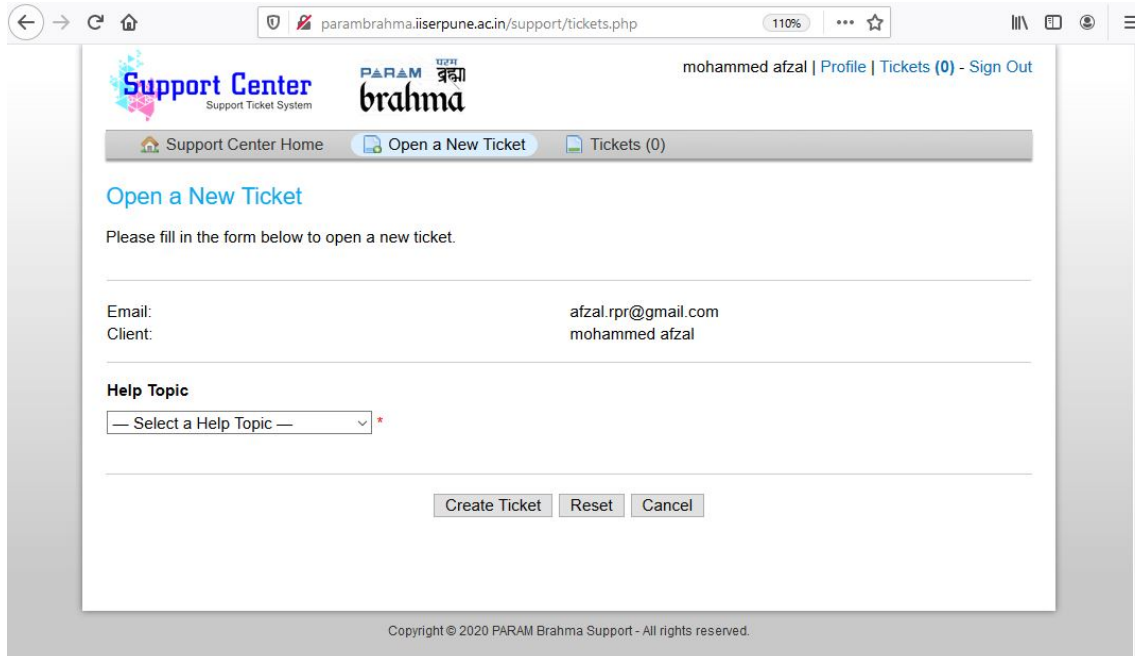


Figure 38 - Snapshot of Ticketing System

5. Please fill in the details of your issue in the fields given and then click on Create ticket.

Figure 39 - Snapshot of Ticketing System

Once the Ticket is generated, an acknowledgement e-mail will be sent to your official e-mail address. The e-mail will also contain the Ticket number along with reference to the ticket that you have generated.


In case of any difficulty while accessing Brahma-Support you can reach us via e-mail at brahmasupport@iiserpune.ac.in.

Closing Your Account on PARAM Brahma

When once you have completed your research work and you no longer need to use PARAM Brahma, you may please close your account on PARAM Brahma. Please raise a ticket by following the url <https://parambrahma.iiserpune.ac.in/support>. The system administrator will guide you about the “Closure Procedure”. You will need clearance from your project-

coordinator/ Supervisor/ Head of the Department about you having surrendered this resource for getting “no dues” certificate from the institute.

PARAM Brahma Account Request Form



PARAM BRAHMA ACCOUNT REQUEST FORM

User Details:

First Name: _____ Last Name: _____

Organization Name: _____

Department: _____

Designation: _____

(Designation: If student, provide the details below)

Roll No.: _____ Course: _____ Academic Year: _____

Official Email address: _____

Office no.: _____ Mobile no.: _____

(If research, provide the details below)

Nature of the Research: _____

Project Details:

Project Name: _____

Nature of the Project: _____

Brief Description of the Project: _____

Project Start Date: _____ Project Duration: _____

Proposed work on PARAM BRAHMA & Requirement of resources:

PARAM Brahma HPC facility usage policies:

1. The Resources provided to you on PARAM Brahma facility should not be used for any commercial purpose i.e. it is restricted for the academic use like research projects, academic projects, NSM projects, NSM approved MSME projects and scientific projects.
2. Sharing your login credentials with some third person will revoke the responsibility of PARAM Brahma administration committee for data theft and your account will also be disabled. The third person will also be held accountable for misusing the PARAM Brahma facility.
3. It is strictly recommended that you should not run jobs on login node and any such incident reported will result in cancellation of the job and any repeat action will result in closure of your account.
4. You will be responsible for informing the PARAM Brahma administration about your project completion, project cancellation and moving or copying data related to your project from PARAM Brahma.
5. You will be solely responsible for keeping your password strong and safe.
6. If found in any engagement or promotion of activities like hacking, reverse-engineering, violating intellectual property rights on or using the PARAM Brahma facility, you will be barred from having account on any Supercomputer setup under the National Supercomputing Mission.
7. The facility is built with least downtime requirement; however, it depends on various factors like Hardware reliability, Power outage, network outage, scheduled maintenance due to which the facility could be unavailable completely/partially. Notification of all scheduled /

unscheduled maintenance will be made known to the users via Website, Email, broadcast message, newsgroups etc.

8. This facility will not be used for any purpose connected with Chemical or Biological or Nuclear weapons or missiles capable of delivering such Weapons.
9. Acknowledging the usage of the facility is mandatory.

If you use supercomputers and services provided under the National Supercomputing Mission, Government of India, please let us know of any published results including Student Thesis, Conference Papers, Journal Papers and patents obtained.

Performa for Acknowledging the usage:

The support and the resources provided by ‘PARAM Brahma Facility’ under the National Supercomputing Mission, Government of India at the Indian Institute of Science Education and Research, Pune are gratefully acknowledged.

Also, please submit the copies of dissertations, reports, reprints and URLs in which “National Supercomputing Mission, Government of India” is acknowledged to:

HoD, HPC Technologies,
Centre for Development of Advanced Computing,
CDAC Innovation Park,
S.N. 34/B/1,
Panchavati, Pashan,
Pune – 411008
Maharashtra

Email: brahmasupport@iiserpune.ac.in

Communication of your achievements using resources provided by National Supercomputing Mission, will help the Mission in measuring outcomes and gauging the future requirements. This will also help in further augmentation of resources at a given site of National Supercomputing Mission.

I acknowledge the above mentioned usage policies & terms and conditions

User’s signature

Recommended/Not Recommended

Signature of HoD/Head of Organization:

Name: _____

Designation: _____

Department: _____

Official Email address: _____

Approving Authority

Approved/Not Approved

Remarks: _____

Name, Signature and seal of approving authority

References

1. <https://lammps.sandia.gov/>
2. <https://www.openacc.org/>
3. <https://www.openmp.org/>
4. <https://computing.llnl.gov/tutorials/mpi/>
5. <https://developer.nvidia.com/cuda-zone>
6. <https://www.mmm.ucar.edu/weather-research-and-forecasting-model>
7. <http://www.gromacs.org/>
8. <https://www.openfoam.com/>
9. <https://slurm.schedmd.com/>
10. https://www.tutorialspoint.com/gnu_debugger/what_is_gdb.htm
11. <https://nsmindia.in/>
12. https://en.wikipedia.org/wiki/Deep_learning
13. <https://docs.conda.io/en/latest/>
14. <https://docs.conda.io/en/latest/miniconda.html>
15. <https://www.tensorflow.org/>
16. <https://www.tensorflow.org/install>
17. <https://github.com/PaddlePaddle/Paddle>
18. Keras, <https://keras.io/>
19. Pytorch, <https://pytorch.org>
20. <https://mxnet.apache.org>
21. <https://software.intel.com/en-us/distribution-for-python>
22. <https://software.intel.com/en-us/articles/intel-optimization-for-tensorflow-installation-guide>